

# Planning, Learning and Decision Making

Lecture 3. Hidden Markov models

# Basic notions



# The urn problem

# The urn problem

- An oracle has available two urns
  - Each urn is filled with black and white balls
  - **Urn 1:** 25% white balls; 75% black balls
  - **Urn 2:** 25% black balls; 75% white balls

# The urn problem

- At each step, the oracle draws a ball from one of the urns
- The ball is put back afterwards

# The urn problem

- The current urn determines which urn the next ball will come from
  - The next ball will come from the same urn with 80% probability
  - The next ball will come from a different urn with 20% probability

# 1. Is this a Markov chain?


- **Yes**
  - The urn-ball in the next step depends on the urn-ball in the current step

## 2. What are the states?

- All possible urn-ball pairs:
  - $\mathcal{X} = \{(1, w), (1, b), (2, w), (2, b)\}$

# 3. Transition probabilities?

- We can represent the game with a transition diagram:




# 3. Transition probabilities?

- We can also represent the transition probabilities as a matrix:


	(1, w)	(1, b)	(2, w)	(2, b)
(1, w)	0.2	0.6	0.15	0.05
(1, b)	0.2	0.6	0.15	0.05
(2, w)	0.05	0.15	0.6	0.2
(2, b)	0.05	0.15	0.6	0.2

# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the next state?




# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
  - What is the next state?
    - $(1, w)$  with probability 0.2
    - $(1, b)$  with probability 0.6
    - $(2, w)$  with probability 0.15
    - $(2, b)$  with probability 0.05
- 
- The diagram illustrates the possible transitions from the state  $(1, b)$ . A vertical red line separates the list of states from the probabilities. To the right of the line, the probability 0.6 is labeled "80% probability" in red, and the probability 0.15 is labeled "20% probability" in red.

# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the state at time  $t = 3$ ?




# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the state at time  $t = 3$ ?
  - $(1, w)$  with probability 0.15
  - $(1, b)$  with probability 0.46
  - $(2, w)$  with probability 0.29
  - $(2, b)$  with probability 0.10

# The urn problem (revisited)

- At each step, the oracle draws a ball from one of the urns
- The ball is put back afterwards
- **The oracle reveals only the color of the ball**

Partial observability



The state cannot be  
fully observed

# 1. Is this a Markov chain?

- **No!**
  - The information available at each moment (ball color) is not enough to predict next ball color
  - Why?
  - The next distribution depends on the urn!

# Can we make predictions?

- Suppose that we know that:
  - At time  $t = 0$ , the ball was white and drawn from urn 1
  - At time  $t = 1$ , the ball was black
  - What is the state at time  $t = 2$ ?

# Can we make predictions?


- We have:

$$t = 0$$

- $x_0 = (1, w)$


# Can we make predictions?

- We have:




# Can we make predictions?

- We have:




# Can we make predictions?

- We have:




# Can we make predictions?

- We have:



# Can we make predictions?

- We have:




# Can we make predictions?

- We have:

$t = 0$

$t = 1$

$t = 2$




# Can we make predictions?

- We have:

$$t = 0$$

$$t = 1$$

$$t = 2$$



# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = x \mid x_0 = (1, w), z_1 = b]$$


$$= \frac{1}{\rho} \left( \mathbf{P}(x \mid (1, b)) \mathbf{P}((1, b) \mid (1, w)) + \mathbf{P}(x \mid (2, b)) \mathbf{P}((2, b) \mid (1, w)) \right)$$


Normalization

Observation

Combination of  
similar terms

First level terms







# Hidden Markov models

# Hidden Markov model

- Model for **sequential process with partial observability**
- Process evolves in **discrete time steps** (as in Markov chains)
- **State of the process** at time step  $t$ :  $x_t$
- **Observation** at time step  $t$ :  $z_t$



# Hidden Markov model



# Hidden Markov model

## Markov state

The state at instant  $t$  is enough to predict the state at instant  $t + 1$ :

$$\mathbb{P} [x_{t+1} = y \mid x_{0:t} = \mathbf{x}_{0:t}, z_{0:t} = \mathbf{z}_{0:t}] = \mathbb{P} [x_{t+1} = y \mid x_t = x_t]$$



Depends only on  
the last state

# Hidden Markov model

## State-dependent observations

The state at instant  $t$  is enough to predict the observation at instant  $t$ :

$$\mathbb{P} [z_t = z \mid \mathbf{x}_{0:t} = \mathbf{x}_{0:t}, \mathbf{z}_{0:t-1} = \mathbf{z}_{0:t-1}] = \mathbb{P} [z_t = z \mid \mathbf{x}_t = x_t]$$



Depends only on  
the last state

# Hidden Markov model

## Markov state

The state at instant  $t$  is enough to predict the state at instant  $t + 1$ :

$$\mathbb{P} [x_{t+1} = y \mid x_{0:t} = \mathbf{x}_{0:t}, z_{0:t} = \mathbf{z}_{0:t}] = \mathbb{P} [x_{t+1} = y \mid x_t = x_t]$$

## State-dependent observations

The state at instant  $t$  is enough to predict the observation at instant  $t$ :

$$\mathbb{P} [z_t = z \mid x_{0:t} = \mathbf{x}_{0:t}, z_{0:t-1} = \mathbf{z}_{0:t-1}] = \mathbb{P} [z_t = z \mid x_t = x_t]$$

# Hidden Markov model

- Other assumptions (for most of this course):
  - There is only a **finite number** of possible states
  - $\mathcal{X}$  is the set of possible states (**state space**)
  - There is only a **finite number** of observations
  - $\mathcal{Z}$  is the set of possible observations (**observation space**)

# Hidden Markov model

- Other assumptions:
  - The transition probabilities  $\mathbb{P} [x_{t+1} = y \mid x_t = x]$  do not depend on  $t$
  - The observation probabilities  $\mathbb{P} [z_t = z \mid x_t = x]$  do not depend on  $t$

# Transition probability matrix


- As with Markov chains, we store the transition probabilities in a **matrix  $P$**

$$[P]_{xy} = \mathbb{P} [x_{t+1} = y \mid x_t = x]$$

# Observation probability matrix

- Similarly, we store the observation probabilities in a **matrix**
  - o

$$[\mathbf{O}]_{xz} = \mathbb{P} [z_t = z \mid x_t = x]$$



Number in row  $x$  column  
 $z$  is the probability of  
observing  $z$  in  $x$

# Example

- **The urn problem**

	(1, w)	(1, b)	(2, w)	(2, b)
(1, w)	0.2	0.6	0.15	0.05
(1, b)	0.2	0.6	0.15	0.05
(2, w)	0.05	0.15	0.6	0.2
(2, b)	0.05	0.15	0.6	0.2

	w	b
(1, w)	1	0
(1, b)	0	1
(2, w)	1	0
(2, b)	0	1

# Summarizing...

- A HMM can be represented compactly as a tuple

$$(\mathcal{X}, \mathcal{Z}, \mathbf{P}, \mathbf{O})$$


- $\mathcal{X}$  is the set of possible states
- $\mathcal{Z}$  is the set of possible observations
- $\mathbf{P}$  is the transition probability matrix
- $\mathbf{O}$  is the observation probability matrix



# Examples


# 1. Household robot

- Remember?




# 1. Household robot

- Robot navigates using a laser



# 1. Household robot

- For example:



# 1. Household robot

- However, laser is not perfect
  - It fails to detect existing walls with 5% probability
  - It detects non-existing walls with 10% probability (in some situations with 20% probability)
  - Detection of a wall independent of adjacent walls

# Is this an HMM?

- State verifies the Markov property?
  - **Yes** - the position of the robot at time  $t + 1$  depends only of position of the robot at time  $t$
- Observation depends only on state?
  - **Yes** - the wall detections depends only on robot position


**It is an HMM!**

# What are the states?

- Possible positions of the robot:
  - $\mathcal{X} = \{K, P, D, H_1, B, L\}$


# What are the observations?

- Possible wall configurations:



# What are the observations?


- Possible wall configurations:
  - $\mathcal{Z} = \{0000, 0001, 0010, 0011,$   
 $0100, 0101, 0110, 0111,$   
 $1000, 1001, 1010, 1011,$   
 $1100, 1101, 1110, 1111\}$



U-D-L-R

# Transition probabilities

- We can represent the underlying state transitions using a transition diagram:




# Transition probabilities

- ... or as a matrix:

$$P = \begin{bmatrix} 0.0 & 0.5 & 0.5 & 0.0 & 0.0 & 0.0 \\ 0.5 & 0.0 & 0.0 & 0.5 & 0.0 & 0.0 \\ 0.5 & 0.0 & 0.0 & 0.5 & 0.0 & 0.0 \\ 0.0 & 0.5 & 0.0 & 0.0 & 0.5 & 0.0 \\ 0.0 & 0.0 & 0.0 & 0.5 & 0.0 & 0.5 \\ 0.0 & 0.0 & 0.0 & 0.5 & 0.5 & 0.0 \end{bmatrix}$$


# Observation probabilities

- The observation probabilities can be computed for each wall independently:



# Observation probabilities

- For example:  $\mathbb{P}[z_t = 1100 \mid x_t = B]$



# Observation probabilities

- This yields a large matrix:

$$O = \begin{bmatrix} 0.0018 & 0.0004 & 0.0342 & \dots & 0.0009 & 0.0722 & 0.0180 \\ 0.0001 & 0.0021 & 0.0021 & \dots & 0.0045 & 0.0045 & 0.0857 \\ 0.0020 & 0.0002 & 0.0384 & \dots & 0.0005 & 0.0812 & 0.0090 \\ 0.5184 & 0.0576 & 0.1296 & \dots & 0.0016 & 0.0036 & 0.0004 \\ 0.0020 & 0.0385 & 0.0002 & \dots & 0.0812 & 0.0004 & 0.0090 \\ 0.0324 & 0.0036 & 0.0081 & \dots & 0.0076 & 0.0171 & 0.0019 \end{bmatrix}.$$

# HMM problems

# HMM problems

- **Estimation:**
  - Estimate HMM trajectories given an HMM model
- ~~Inference:~~
  - Determine an HMM model from data

# Estimation

- **Filtering:** Forward algorithm
  - Given a sequence of observations, estimate the final state
- **Smoothing:** Viterbi algorithm
  - Given a sequence of observations, estimate the sequence of states
- **Prediction:**
  - Given a sequence of observations, predict future states



# Filtering

# Filtering

- We are given a sequence of observations  $\mathbf{z}_{0:T}$
- We want to estimate

$$\mathbb{P}_{\mu_0} [\mathbf{x}_T = \mathbf{x} \mid \mathbf{z}_{0:t} = \mathbf{z}_{0:T}]$$

where  $\mu_0$  is the initial distribution, i.e.,

$$\mu_0(\mathbf{x}) = \mathbb{P} [\mathbf{x}_0 = \mathbf{x}]$$

# Filtering

- For simplicity, write

$$\mu_{T|0:T}(x) = \mathbb{P}_{\mu_0} [\mathbf{x}_T = x \mid \mathbf{z}_{0:t} = \mathbf{z}_{0:T}]$$

# Forward mapping

## Forward mapping

Given a sequence of observations  $\mathbf{z}_{0:t}$ , the forward mapping  $a_t : \mathcal{X} \mapsto \mathbb{R}$  is defined for each  $t$  as

$$\alpha_t(x) = \mathbb{P}_{\mu_0} [\mathbf{x}_t = x, \mathbf{z}_{0:t} = \mathbf{z}_{0:t}]$$

# So what?

- Forward mapping has several useful properties
  1. We can compute  $\mu_{T|0:T}$  from  $a_T$ :

$$\mu_{T|0:T}(x) = \frac{\alpha_T(x)}{\sum_{y \in \mathcal{X}} \alpha_T(y)}$$

# So what?

- Forward mapping has several useful properties
  1. We can compute  $\mu_{T|0:T}$  from  $a_T$
  2. The forward mapping can be computed recursively:

$$\alpha_T(x) = \mathbf{O}(z_T \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{T-1}(y)$$

# Forward algorithm

**Require:** Observation sequence  $z_0:T$

1. Initialize  $\alpha_0 \leftarrow \text{diag}(\mathbf{O}_{:,z_0})\boldsymbol{\mu}_0^\top$

2. **for**  $t = 1, \dots, T$  **do**

$$\alpha_t \leftarrow \text{diag}(\mathbf{O}_{:,z_t})\mathbf{P}^\top \alpha_{t-1}$$

If we have an  
initial observation

**4. end for**

5. **return**  $\boldsymbol{\mu}_{T|0:T} = \alpha_T / (\mathbf{1}^\top \alpha_T)$

# Example: The urn problem

- Suppose that

$$\mu_0 = [0.125 \quad 0.375 \quad 0.375 \quad 0.125]$$


- We observe the sequence of observations

$$z_{0:2} = \{w, w, b\}$$

- What is the state at time  $t = 2$ ?


# Step 1: Initialize $a_0$

- $\alpha_0 \leftarrow \text{diag}(\mathbf{O}_{:,z_0})\boldsymbol{\mu}_0^\top$



# Step 2: Compute $\alpha_1$

- $\alpha_1 \leftarrow \text{diag}(\mathbf{O}_{:,z_1})\mathbf{P}^\top\alpha_0$


$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.2 & 0.2 & 0.05 & 0.05 \\ 0.6 & 0.6 & 0.15 & 0.15 \\ 0.15 & 0.15 & 0.6 & 0.6 \\ 0.05 & 0.05 & 0.2 & 0.2 \end{bmatrix} \begin{bmatrix} 0.125 \\ 0 \\ 0.375 \\ 0 \end{bmatrix} \xrightarrow{\text{red arrow}} \begin{bmatrix} 0.04 \\ 0 \\ 0.24 \\ 0 \end{bmatrix}$$

# Step 3: Compute $\alpha_2$

- $\alpha_2 \leftarrow \text{diag}(\mathbf{O}_{:,z_2})\mathbf{P}^\top\alpha_1$



# Final step: Compute $\mu_{2|0:2}$

- We finally get:

$$\begin{aligned}\mu_{2|0:2} &= \alpha_2 / (\mathbf{1}^\top \alpha_2) \\ &= [0 \quad 0.552 \quad 0 \quad 0.448]\end{aligned}$$