

Semantics



Definition of Security Properties

Ana Matos

Miguel Correia


Pedro Adão

“Secure” programs

- We want to ensure that programs are “secure”.
- First, we must be clear about what we are protecting, what is the power of the attacker, and what being secure actually means.

Example: E-Voting


- Functionality
 - Cast vote, audit vote, determine winner, ...
- Assets
 - votes, voting service,
- Threat model
 - Goal: the adversary desires to influence the election.
 - Methods: Voting coercion, disabling or controlling voting booths, communications, servers...
- (...)



- (...)
- Security properties
 - Every citizen can cast a vote. (availability)
 - The identity of who placed each vote is anonymous, and the choice of each voter is known only to the voter. (confidentiality)
 - Only the voter can determine its vote. (integrity)
 - The voter can verify that its vote was casted correctly. (auditability)...
- Enforcement mechanism
 - Authentication, cryptographic protocols ensure (up to cryptographic strength) anonymity, integrity and auditability, ...

How can we get strong guarantees?

We need to be precise.

- 
- “Is the program secure?”
 - “Is the program secure against this threat?”
 - “Is the program secure against attacker model A?”
 - “Does the program meet security property P in the execution context C?”
 - “Can we prove that the program meets security property P in the execution context C?”

Formal security property

- Why use a **formal** security property?
 - For clarity (preventing ambiguities in the specification)
 - For conciseness
 - For correctness (the basis for implementation, analysis and verification)
- What do we need?
 - Techniques for reasoning about the syntax and the semantics of programs.

Class Outline

- Noninterference, intuitively
- Formal semantics
 - Big-step operational semantics
- Formalization of Noninterference

Class Outline

- Noninterference, intuitively
- Formal semantics
 - Big-step operational semantics
- Formalization of Noninterference

Noninterference, intuitively



“Language-Based Information-Flow Security”, A. Sabelfeld and A. Myers , 2002.

The property of being “secure”

- We want to ensure that propagation of information by programs respects information flow policies: an attacker cannot infer secret input or affect critical output by inserting inputs into the system and observing its outputs.
- How can we express the property of whether a program respects information flow policies?

Suppose YOU are the attacker


- You can only read and write to variables of “low level” -- i.e., your level or lower
- Can you use the program to uncover “high level” information?
- Can you use the program to affect “high level” information?

Challenge for next class

- Secure program = the program does not encode illegal information flows
- Can you formulate a security property that defines when a program is secure?

Tips

- Suppose you are the attacker. You are given a black box program, and you can control the low inputs of the program, and read its low outputs.
- How many executions would you need to find out if the program leaks information?
- What inputs would you give it in order to find out whether the program leaks information?




- Which of the following experiments allows you to conclude whether the program encodes leaks?
- Input: $y_L=0$, Output: $y_L=0$
- Input: $y_L=0$, Output: $y_L=1$
- Input: $y_L=0$, Output: $y_L=0$ and Input: $y_L=0$, Output: $y_L=0$
- Input: $y_L=0$, Output: $y_L=0$ and Input: $y_L=0$, Output: $y_L=1$

What is the power of the attacker?

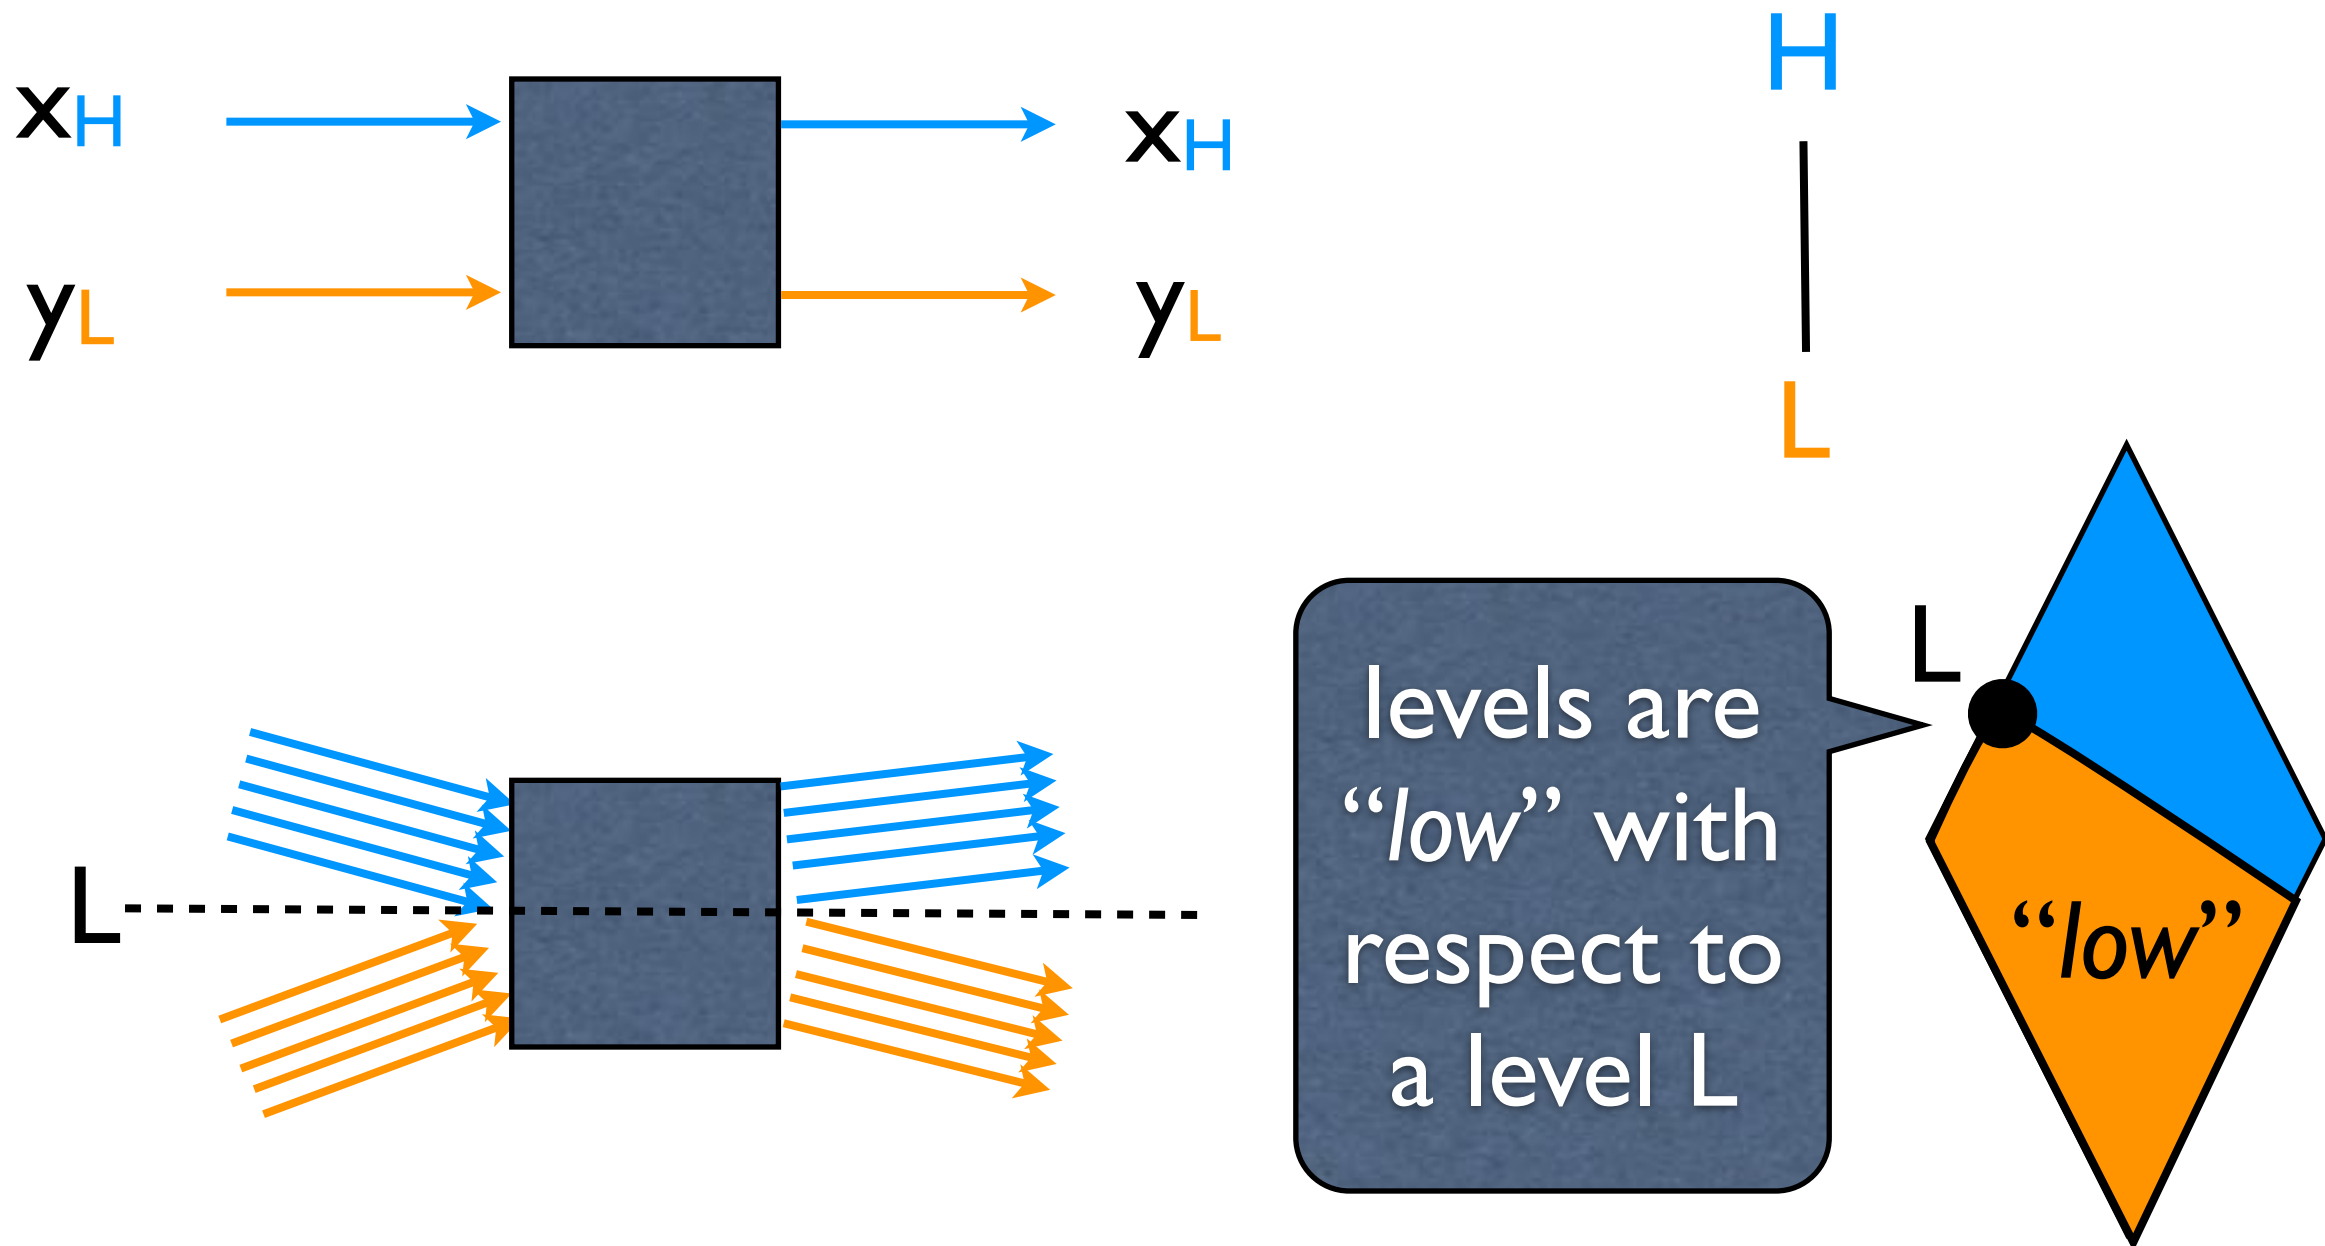
- Execution context: sequential (deterministic)
- Intuition: an attacker is a program that is sequentially composed with S , and has access to “low” outputs.
- Deterministic Input-Output attacker:
Can only observe final outputs of terminating computations.

Noninterference (almost)




- The program is secure if, for any two runs of the program that are given the same *low* (L) inputs, if the program terminates, it produces the same low (L) outputs.

Noninterference, for general policies



Noninterference (informally)



- A program is secure if, for any observational level L , for any two runs of the program that are given the same *low* inputs, if the program terminates, it produces the same *low* outputs.

Secure? (w.r.t. ...)

- $y_H := x_L$ ✓

- $x_L := y_H$ ✗

Explicit leak

- if y_H then $x_L := 0$ else $x_L := 1$ ✗

- while y_H do skip ; $x_L := 0$ ✓

Implicit leak

Covert information flow

- Consider two files:
 - “secret.txt” - owned by Alice, can only be read and written by Alice
 - “public.txt” - can be read and written by everyone
- Can you conceive a program that discloses the secret information **but respects Deterministic Input-Output Noninterference?**
- Can Discretionary Access Control prevent it?
- Can Mandatory Access Control prevent it?

Limitations of the attacker

- A deterministic input-output attacker model assumes that the attacker cannot observe
 - non-terminating computations
 - intermediate steps of computations
 - time that it takes to produce the output
 - likelihood of each possible output
- This can be inadequate for different
 - execution contexts
 - language expressivity

Concurrent execution context

- Is our notion of Noninterference adequate in a concurrent context?
- $x_L := 1 \parallel (x_L := 2; x_L := x_L + 2)$
- The above program is considered insecure because of its non-deterministic behavior
- $\langle x_L := 1 \parallel (x_L := 2; x_L := x_L + 2), \rho \rangle \rightarrow \rho[x_L := 4]$
- $\langle x_L := 1 \parallel (x_L := 2; x_L := x_L + 2), \rho \rangle \rightarrow \rho[x_L := 1]$
- $\langle x_L := 1 \parallel (x_L := 2; x_L := x_L + 2), \rho \rangle \rightarrow \rho[x_L := 3]$

Concurrent execution context

- The following programs are considered secure with respect to our notion of Noninterference
 - if PIN_H then $y_H := \text{false}$ else $z_H := \text{false}$
 - while y_H do skip ; $x_L := 1$; $z_H := \text{false}$
 - while z_H do skip ; $x_L := 0$; $y_H := \text{false}$
- But when composed concurrently, the program is insecure!

What is the power of the attacker?

- Execution context: concurrent
- Intuition: an attacker program that is concurrently composed with S , does not depend on termination of the observed program. It has access to “low” outputs, and possibly non-termination (or even intermediate steps).
- Possibilistic Input-Output attacker:
Can observe whether the program is capable of producing certain final outputs.

Secure? (w.r.t. ...)

- $y_H := x_L$ ✓
- $x_L := y_H$ ✗
- if y_H then $x_L := 0$ else $x_L := 1$ ✗
- while y_H do skip ; $x_L := 0$ ✗

Termination leak

Counter-example:
 $\rho_1(x_L) = \rho_2(x_L) = 1$ and $\rho_1(y_H) = \text{false}$ and $\rho_2(y_H) = \text{true}$
(the program cannot terminate on ρ_2)

Possibilistic Input-Output Noninterference

- Program S is secure if, for any two memories that agree on “low” variables, if running S on one of them terminates and produces a certain final memory, then running S on the other **can also** terminate and produce a final memory that agrees on the “low” variables.

Termination-sensitive!

Intermediate-step attacker

$x_L := y_H ; x_L := 1$

Possible low outcomes do not depend on y_H .
However, the intermediate steps differ.

- Intermediate-step-sensitive Noninterference:
Is sensitive to intermediate steps of computations.
- To define an intermediate-step-sensitive property
we would use a small-step semantics.

Time-sensitive attacker

$x_L := 0$; if y_H then skip else skip;skip;skip;skip ; $x_L := 1$

Possible outcomes and intermediate steps do not depend on y_H . However, the time it takes to change the value of x_L is different.

- Temporal Noninterference:
Is sensitive to the time it takes to produce outputs.

Probabilistic attacker

$x_L := y_H \parallel x_L := \text{random}(100)$

Possible outcomes do not depend on y_H . However, the probability of the value of x_L revealing that of y_H is higher.

- Probabilistic Noninterference:
Can observe the likelihood of outputs.

Limits of Noninterference

- As sensitive as the attacker model.
 - Covert channels?
- As flexible as what the security property permits.
 - Too restrictive?

Covert channels

- Information can be transferred via other side-channels that are not intended for that purpose:
 - computation time
 - memory allocation
 - power consumption
 - ...

More flexibility

- Noninterference is appealing because it provides strong security guarantees. However it is too restrictive!
- `if (passwordH == attemptL)
 then printL "Right!"
 else printL "Wrong!"`
- `$filename=<STDIN>;
open(FOO,"> $filename");`

Downgrading

- Sometimes we need to leak information in a controlled way.
- Example for confidentiality (called declassification):
 declassify password:L in
 if (password_H == attempt_L)
 then print_L “Right!”
 else print_L “Wrong!”
- And for integrity (called endorsement)?
 Think of Perl’s taint mode.

Conclusion

- The framework for the analysis should reflect the assumptions we want to study:
 - Language expressivity?
 - Execution context?
 - Attacker power?
 - Strength/flexibility of the security requirements?

In this course

- We will focus primarily on Input-Output Deterministic Noninterference.
- We will study different enforcement mechanisms for this security property.
- We want to have strong guarantees of that our mechanisms work.
- We will therefore use formal methods.


To start with, we need

- A precise way to talk about:
 - what are the possible programs and inputs
 - what is the (final) result (output) of executing a (deterministic) program with a certain input.
- In other words, we need:
 - a formal language (syntax and semantics)
 - notation to reason about terminating computations.

Class Outline

- Noninterference, intuitively
- Formal semantics
 - Big-step operational semantics
- Formalization of Noninterference

Semantics?



Programming language semantics

specification of the meaning of programs of
that programming language



“Semantics with Applications: A Formal Introduction”, H. Nielson,
F. Nielson, 1992 (Chapter 1 -- Section 2.1.).

Defining the semantics of a programming language

We will use two techniques:

- Denotational semantics - defines what is the result of a computation, as a mathematical object.
- Operational semantics - describes how the effect of a computation is produced when executed on a machine.

A WHILE language

- The next slides present a WHILE language as defined in the tutorial G. Barthe et al. (2011).
- The language is a (simpler) variation of the one in Nielson & Nielson (1999).(*)

(*) Differences are:

- We refer to the natural operational semantics as “big-step semantics”, and to the structural operational semantics as “small-step semantics”.
- Constants are simply integers (and not numerals).
- Arithmetic expressions use operations that are used directly in the semantics.
- Boolean expressions are reduced to comparisons between arithmetic expressions (tests).

Syntax of WHILE

- Syntactic categories:
 $c \in \mathbb{Z}$ - constants (integers)
 $x \in \text{Var}$ - variables
 $a \in \text{Aexp}$ - arithmetic expressions
 $t \in \text{Bexp}$ - tests
 $S \in \text{Stm}$ - statements
- Grammar (BNF notation):
 $\text{op} ::= + \mid - \mid \times \mid /$ $\text{cmp} ::= < \mid \leq \mid = \mid \neq \mid \geq \mid >$
 $a ::= c \mid x \mid a_1 \text{ op } a_2$ $t ::= a_1 \text{ cmp } a_2$
 $S ::= x := a \mid \text{skip} \mid S_1 ; S_2 \mid \text{if } t \text{ then } S \text{ else } S \mid \text{while } t \text{ do } S$

Semantic functions

- \mathcal{A} - function that maps pairs of arithmetic expression and state to integers (assume defined)
- \mathcal{B} - function that maps pairs of test and state, to booleans (assume defined)
- \mathcal{S} - partial function that maps pairs of statement and state to state (to define next).

State function ρ

What is a state?

- ρ - memory or state function that maps variables to integers

Example: $\rho(x) = 1$, $\rho(y) = 42 \dots$

Semantics of statements: S

- S - partial function that maps pairs (statement, state) to state.
- We will define it using a “big-step” operational semantics, by means of **big-step transitions**: speaks about how the overall result of executions is obtained.

$$\langle S, \rho \rangle \rightarrow \rho'$$

When executing program S on memory ρ we obtain the new memory ρ'

- We could also have used a “small-step” operational semantics, by means of **small-step transitions**: speaks about how the individual steps of the computations take place.

$$\langle S, \rho \rangle \Rightarrow \langle S', \rho' \rangle$$

Performing one step of program S on memory ρ leaves the continuation S' and produces new memory ρ'

- We will use small step semantics later in the course.

Operational semantics (big-step transitions)

- **Skip:** $\langle \text{skip}, \rho \rangle \rightarrow \rho$

Does nothing!

- **Assignment:** $\langle x := a, \rho \rangle \rightarrow \rho[x \mapsto \mathcal{A}[a]\rho]$

Updates the value of x in ρ with the result of evaluating a

the update of state ρ
is defined as:

$$\begin{cases} (\rho[y \mapsto c])(x) = c, & \text{if } x=y \\ \rho(x), & \text{otherwise} \end{cases}$$

Axioms - do not depend on any hypothesis in order to give the final result of the **entire computation**

When the first program
starting on ρ
produces ρ' ...

... and the second program
starting on ρ'
produces ρ'' ...

- **Sequential composition:**

$\langle S_1, \rho \rangle \rightarrow \rho' \quad \langle S_2, \rho' \rangle \rightarrow \rho''$

$\langle S_1; S_2, \rho \rangle \rightarrow \rho''$

... then the entire sequential
composition starting on ρ
produces ρ'' .

Rules - the final result of the **entire computation** below the
line, depends on the hypothesis above the line

When t evaluates
to true...

... and the first branch
starting on ρ produces
 ρ' ...

- Conditional test:

$$\frac{\langle S_1, \rho \rangle \rightarrow \rho'}{\langle \text{if } t \text{ then } S_1 \text{ else } S_2, \rho \rangle \rightarrow \rho'}$$

if $\mathcal{B}\llbracket t \rrbracket_\rho = \text{true}$

then the conditional
starting on ρ produces
 ρ' .

$$\frac{\langle S_2, \rho \rangle \rightarrow \rho'}{\langle \text{if } t \text{ then } S_1 \text{ else } S_2, \rho \rangle \rightarrow \rho'}$$

if $\mathcal{B}\llbracket t \rrbracket_\rho = \text{false}$

When t evaluates to true...

- **While loop:**

$\langle S, \rho \rangle \rightarrow \rho' \quad \langle \text{while } t \text{ do } S, \rho' \rangle \rightarrow \rho''$ **if** $\mathcal{B}\llbracket t \rrbracket_{\rho} = \text{true}$

$\langle \text{while } t \text{ do } S, \rho \rangle \rightarrow \rho''$

... the body starting on ρ produces ρ' ...

... and the continuation of the cycle on ρ' produces ρ'' ...

... then the cycle on ρ produces ρ'' .

$\langle \text{while } t \text{ do } S, \rho \rangle \rightarrow \rho$ **if** $\mathcal{B}\llbracket t \rrbracket_{\rho} = \text{false}$

When t evaluates to false the cycle does nothing.

(All) Big-step axioms & rules

Assignment: $\langle x := a, \rho \rangle \rightarrow \rho[x \mapsto \mathcal{A}[a]_\rho]$

Skip: $\langle \text{skip}, \rho \rangle \rightarrow \rho$

Sequential composition:
$$\frac{\langle S_1, \rho \rangle \rightarrow \rho' \quad \langle S_2, \rho' \rangle \rightarrow \rho''}{\langle S_1; S_2, \rho \rangle \rightarrow \rho''}$$

Conditional test:
$$\frac{\langle S_1, \rho \rangle \rightarrow \rho'}{\langle \text{if } t \text{ then } S_1 \text{ else } S_2, \rho \rangle \rightarrow \rho'} \quad \text{if } \mathcal{B}[t]_\rho = \text{true}$$
$$\frac{\langle S_2, \rho \rangle \rightarrow \rho'}{\langle \text{if } t \text{ then } S_1 \text{ else } S_2, \rho \rangle \rightarrow \rho'} \quad \text{if } \mathcal{B}[t]_\rho = \text{false}$$

While loop:
$$\frac{\langle S, \rho \rangle \rightarrow \rho' \quad \langle \text{while } t \text{ do } S, \rho' \rangle \rightarrow \rho''}{\langle \text{while } t \text{ do } S, \rho \rangle \rightarrow \rho''} \quad \text{if } \mathcal{B}[t]_\rho = \text{true}$$
$$\langle \text{while } t \text{ do } S, \rho \rangle \rightarrow \rho \quad \text{if } \mathcal{B}[t]_\rho = \text{false}$$

Example - Evaluation

- Evaluate $(z:=x ; x:=y) ; y:=z$, starting from a state ρ_0 that maps all variables except x and y to 0, and has $\rho_0(x) = 5$ and $\rho_0(y) = 7$.

$$\langle z:=x, \rho_0 \rangle \rightarrow \rho_1$$

$$\langle x:=y, \rho_1 \rangle \rightarrow \rho_2$$

derivation tree

$$\langle (z:=x ; x:=y), \rho_0 \rangle \rightarrow \rho_2$$

$$\langle y:=z, \rho_2 \rangle \rightarrow \rho_3$$

$$\langle (z:=x ; x:=y) ; y:=z, \rho_0 \rangle \rightarrow \rho_3$$

$$\rho_1 = \rho_0[z \mapsto 5] \quad \rho_2 = \rho_1[x \mapsto 7] \quad \rho_3 = \rho_2[y \mapsto 5]$$

Evaluation - derivation tree

To evaluate a statement S , starting from a state ρ , a **derivation tree** must be constructed:

1. Construct the tree from root upwards.
2. Try to find an axiom or rule whose left side matches $\langle S, \rho \rangle$ and whose side conditions are satisfied.
 - If it is an axiom - determine the final state and terminate.
 - If it is a rule - try to construct the derivation tree for the premisses of the rule in order to determine the final state.

The semantic function S

- The meaning of statements is given as a (partial) function from the set of states to the set of states.

$$S \llbracket S \rrbracket_{\rho} = \begin{cases} \rho' & \text{if } \langle S, \rho \rangle \rightarrow \rho' \\ \text{undefined,} & \text{otherwise} \end{cases}$$

no meaning is given to
non-terminating
computations

Summarizing: Big-step transition system

- Configurations:
 - intermediate $\langle \text{Statement } S, \text{state } \rho \rangle$
 - terminal ρ
- Transitions: $\langle S, \rho \rangle \rightarrow \rho'$
- Rules:
$$\frac{\langle S_1, \rho_1 \rangle \rightarrow \rho_1' \dots \langle S_n, \rho_n \rangle \rightarrow \rho_n'}{\langle S, \rho \rangle \rightarrow \rho'} \quad \text{if } \dots$$

Conclusions

- We have defined formally the operational semantics for a simple language WHILE.
- The same can be done **for any language in a similar** manner (in future class, for a byte-code language.)
- This allows us to **reason rigorously about the behavior** of programs in the language.
- We will use it to **formalize and prove security properties**, which are often semantic.

Class Outline

- Noninterference, intuitively
- Formal semantics
 - Big-step operational semantics
- Formalization of Noninterference

Formalization of Noninterference



“Language-Based Information-Flow Security”, A. Sabelfeld and A. Myers , 2002.

Noninterference (informally)

Noninterference -

A program is secure if, for any observational level L , for any two runs of the program that are given the same *low* inputs, if the program terminates, it produces the same *low* outputs.

Noninterference (informally)

Program from a given programming language.

Given an information flow policy, all security levels that are lower or equal to L .

Noninterference -


A **program** is secure if, for any **observational level L** , for any two **runs** of the program that are given the **same low input**, if the program terminates, it produces the **same low outputs**.

Execution of the program according to specified semantics.

Security labeling and indistinguishability relation between states.

Indistinguishability

Two memories ρ_1 and ρ_2 are **indistinguishable** with respect to a security labeling Γ and a level **L**, if ρ_1 and ρ_2 agree on the values of variables that are lower or equal to L. I.e.: For all variables x such that $\Gamma(x) \leq L$ we have that $\rho_1(x) = \rho_2(x)$.



We then write $\rho_1 \sim_L \rho_2$.

(Omitting the parameter Γ for simplicity.)

Noninterference (formally)

Deterministic Input-Output Noninterference -

Program S is secure if for every security level L
and for all pairs of memories ρ_1 and ρ_2 such that

$\rho_1 \sim_L \rho_2$, we have that

$\langle S, \rho_1 \rangle \rightarrow \rho_1'$ and $\langle S, \rho_2 \rangle \rightarrow \rho_2'$ implies $\rho_1' \sim_L \rho_2'$.

Conclusions

- In order to have strong guarantees about security, we need to **be precise** about the security property we want to enforce.
- Defining a security property requires defining what is the **threat model**, including the execution context and the power of the attacker.
- To be precise about a (semantic) security property requires a **formal semantics** for expressing it.

