


# Introduction to Information Flow Security

Ana Matos

Pedro Adão


Miguel Correia

# Who should access what?



# “End-to-end” policies

- Confidentiality / Integrity  
Information of any level should only be readable / affected by users of levels that are the same or higher.
- Policies are respected throughout the entire process, as information flows through the system.



# To think before next class

- Can we prevent illegal information flows by means of access control?
- How could we do better?

# Class Outline

- Information Flow Security
  - introduction
  - Access Control to Information Flow Control
- Encoding and exploiting information flows


# Access Control to Information Flow Control

☞ “Lattice-Based Access Control Models”, R. Sandhu, 1993.

☞ “Language-Based Information-Flow Security”, A. Sabelfeld and A. Myers, 2002.

# Access control

- The control of interaction between subjects (active entities such as processes, ...) and objects (protected entities such as files, ...), by **validating access rights** of subjects to resources of the system.



- **Decisions do not take into account how information is propagated by programs.**

# End-to-end policies

- Consider three files:
  - secret.txt (only Alice is allowed to read)
  - crucial.txt (only Bob is allowed to write)
  - public.txt, (both are allowed to read and write)
- Should Alice be allowed to copy the contents of secret.txt to public.txt? Should Bob be allowed to copy the contents of public.txt to crucial.txt?
- Can access control prevent it from happening?



# Discretionary Access Control (DAC)

- Restricts access based on the identity of subjects and a set of access permissions that **can be determined by subjects**.
- Limitations: Access permissions might allow programs to, in effect, circumvent the policies:
  - Legally, by means of information flows that are encoded in the program.
  - Illegally, when vulnerabilities in programs and language implementations (buffer overflows, race conditions,...) can be exploited by attackers.

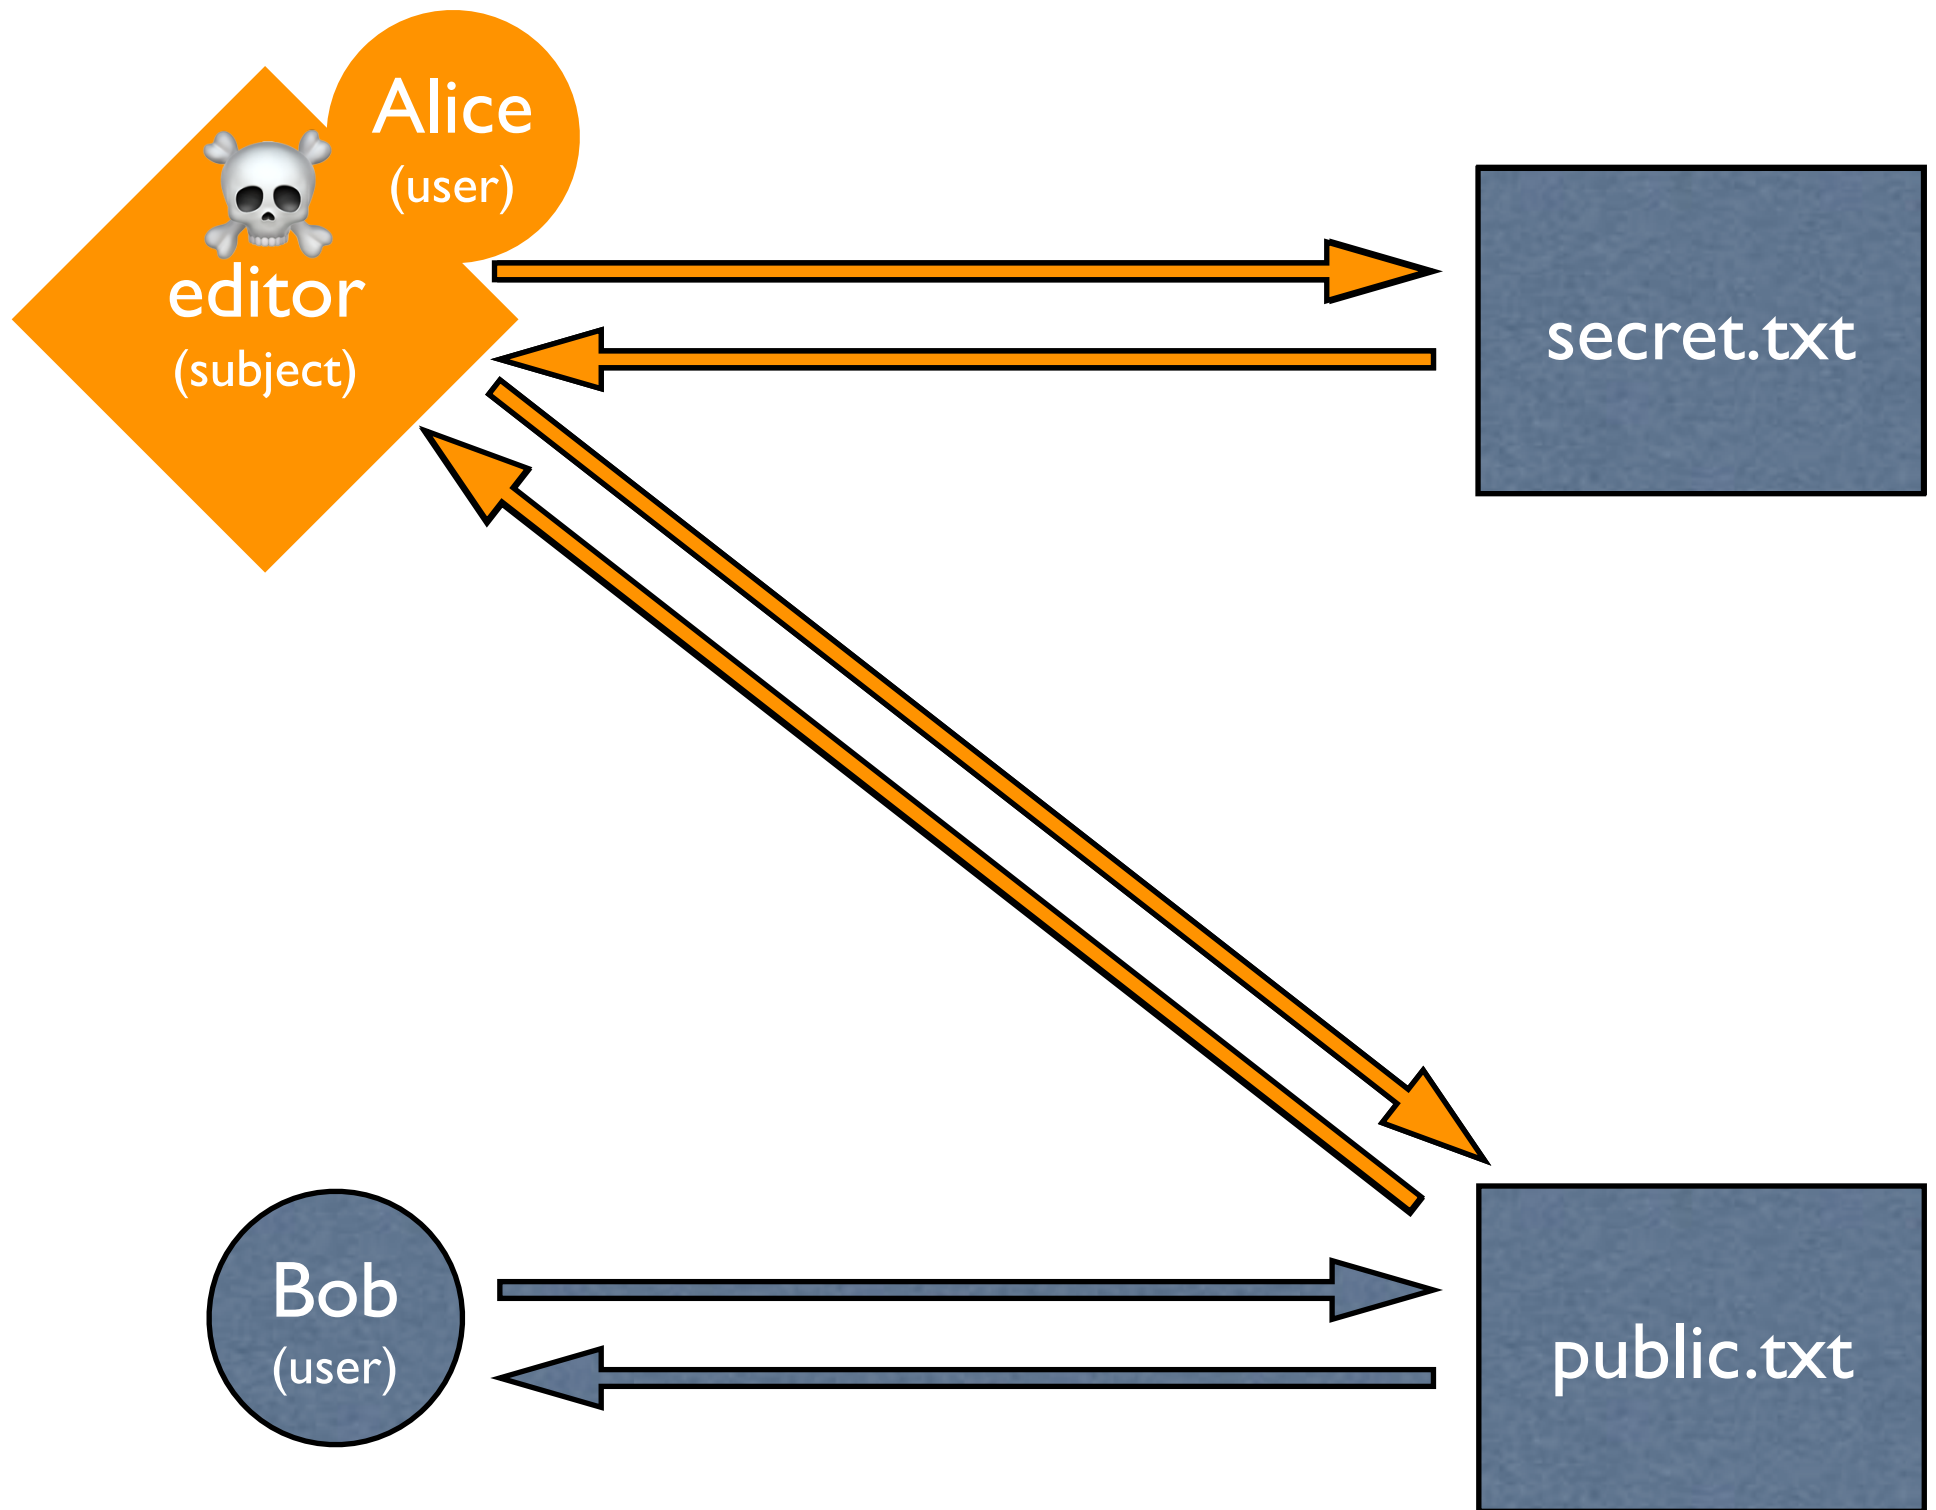
# Example

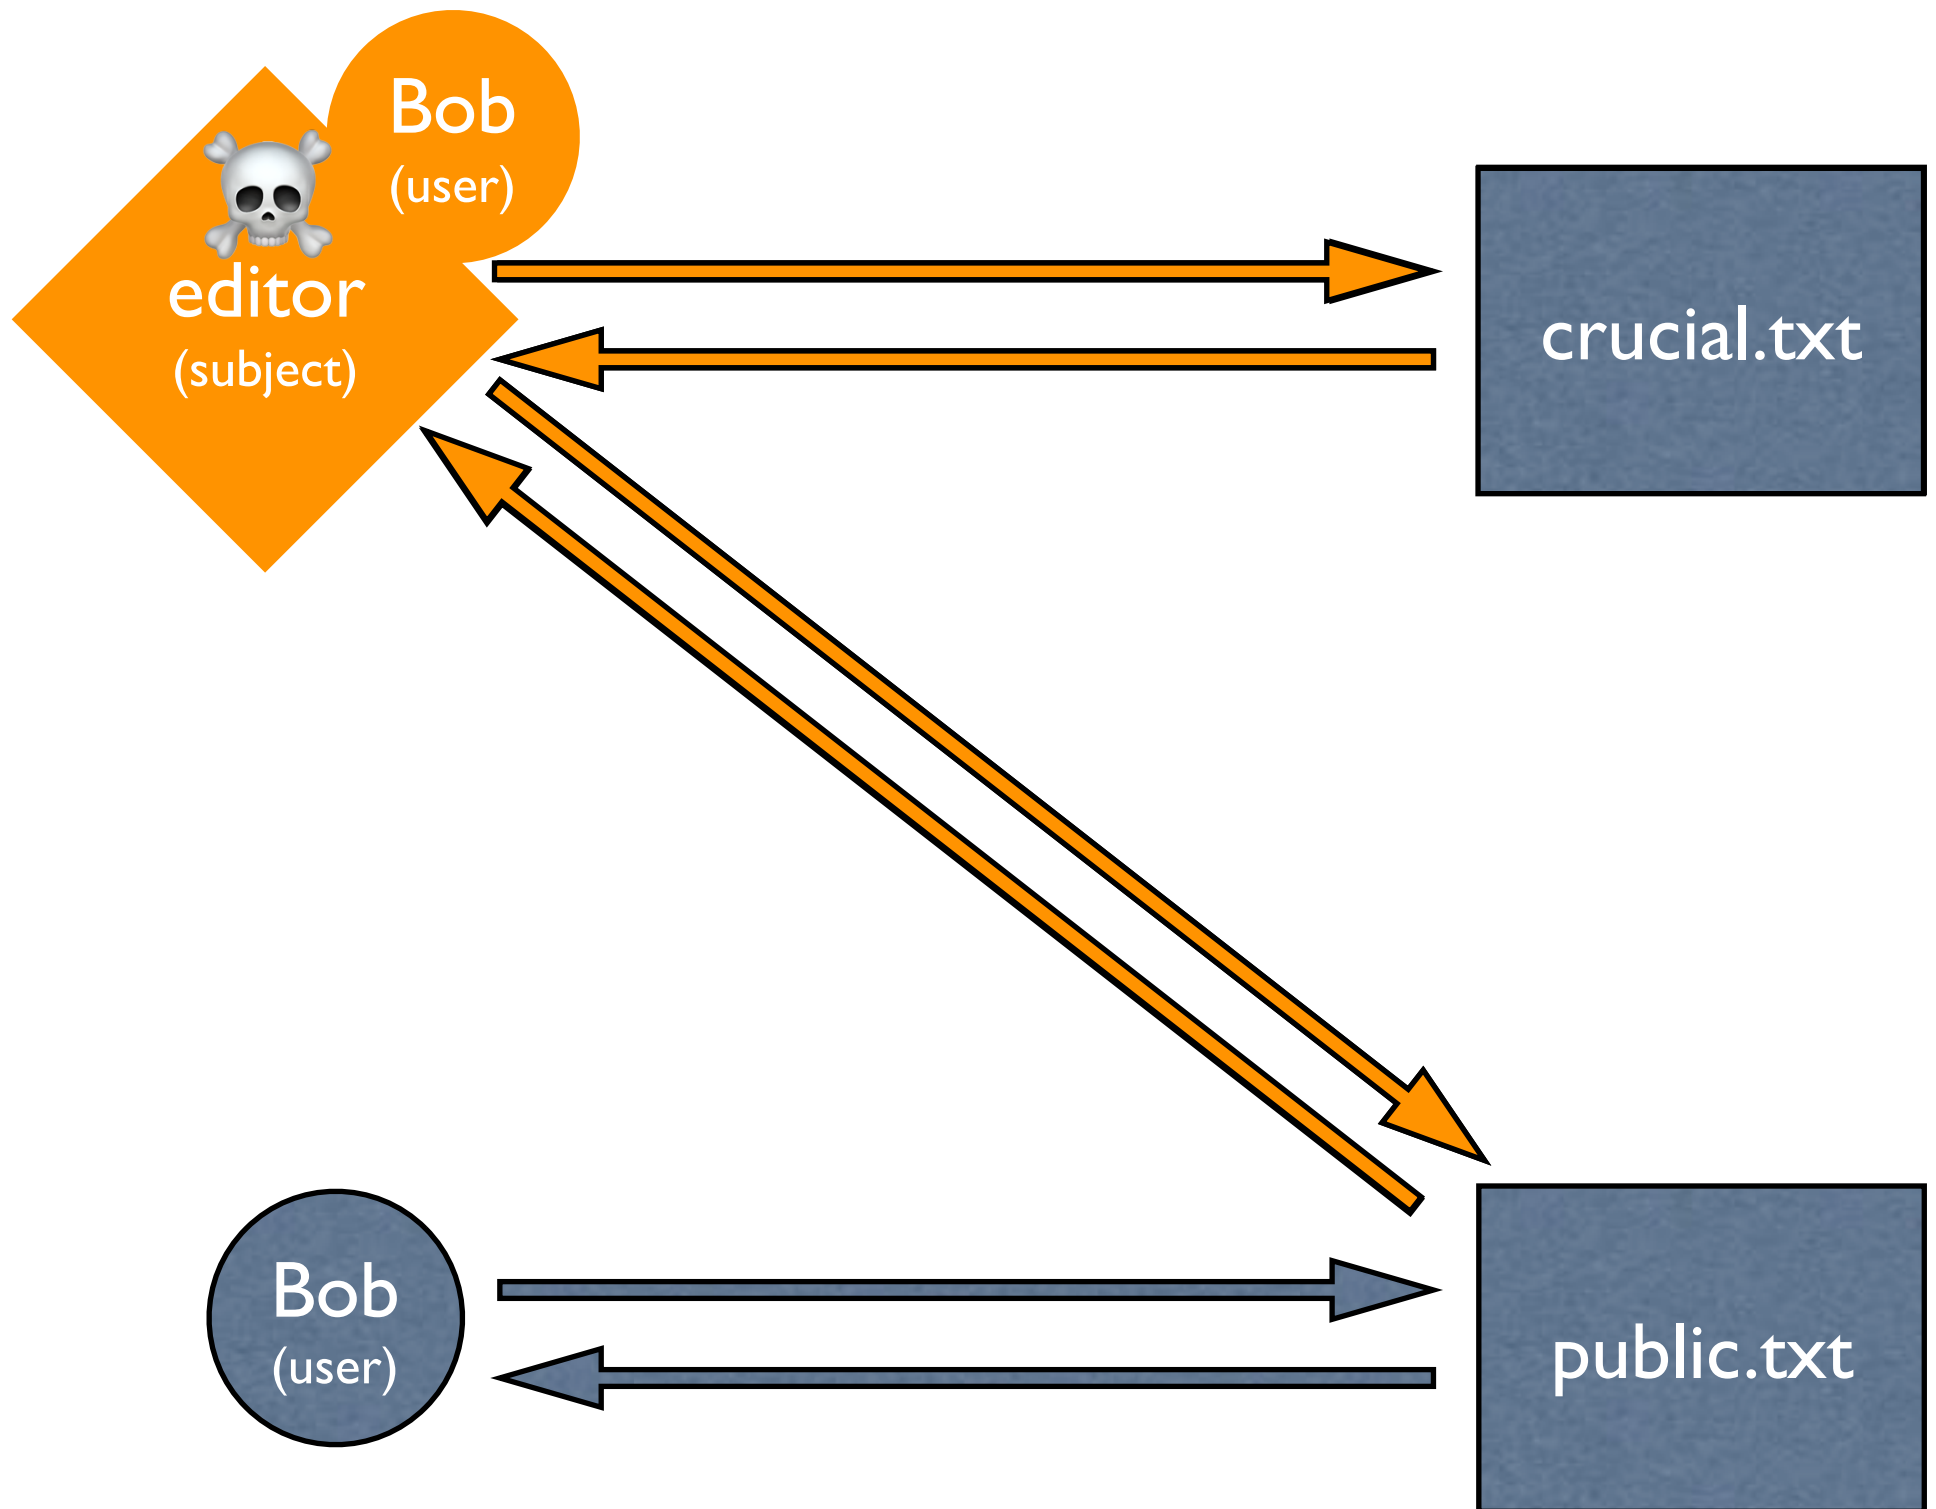
Authority of the process	secret.txt	crucial.txt	public.txt
Alice	r	-	r, w
Bob	-	w	r, w

- Can you conceive a program that discloses Alice's secret information?  
Or ruins Bob's crucial information?
- Which owner and permissions would it have?
- Would it require Alice's (resp. Bob's) cooperation?

# Trojan Horse



- Access is validated or not depending only on the current authority and access permissions.
- Users do not always know what the program they are executing does. A program can perform unwanted things under the user's permissions.
- If a program is exploitable, the attacker can make use of its EUID (root?) privileges
- Out of reach of DAC.





# Mandatory Access Control (MAC)

- Restricts access based on **security levels** of subjects (their clearances) and objects (their sensitivity). Controls how information flows in a system **based on whom** is performing each access.




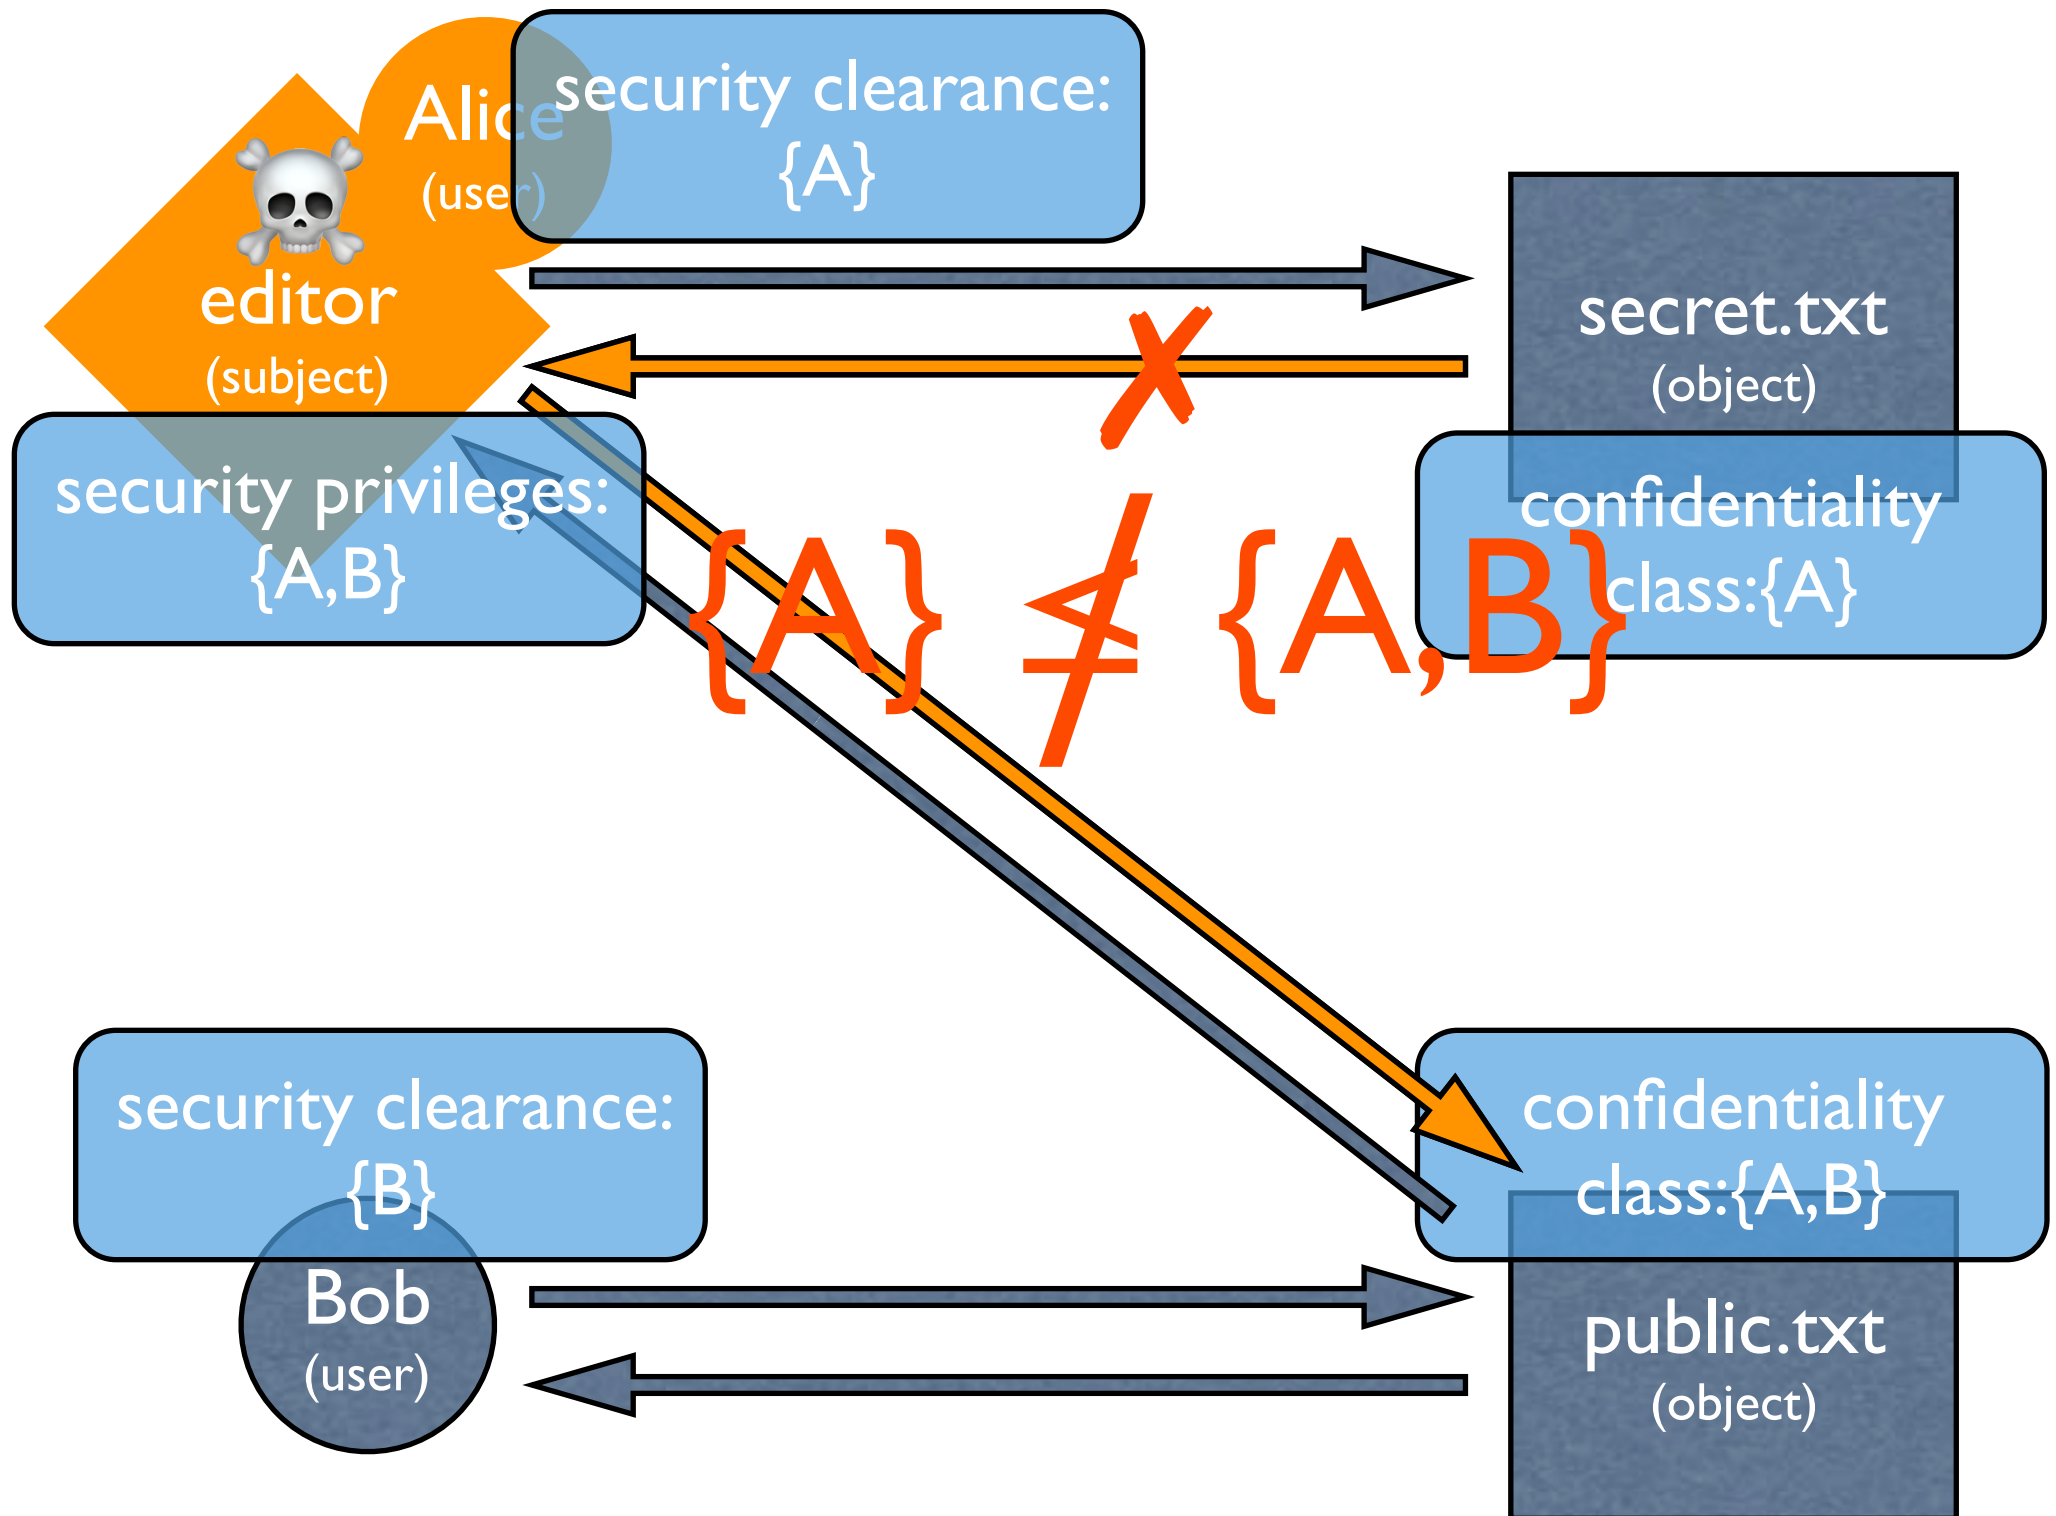
Example: The military hierarchy TS - S  
- C - U.

Limitations: restrictiveness, covert channels.

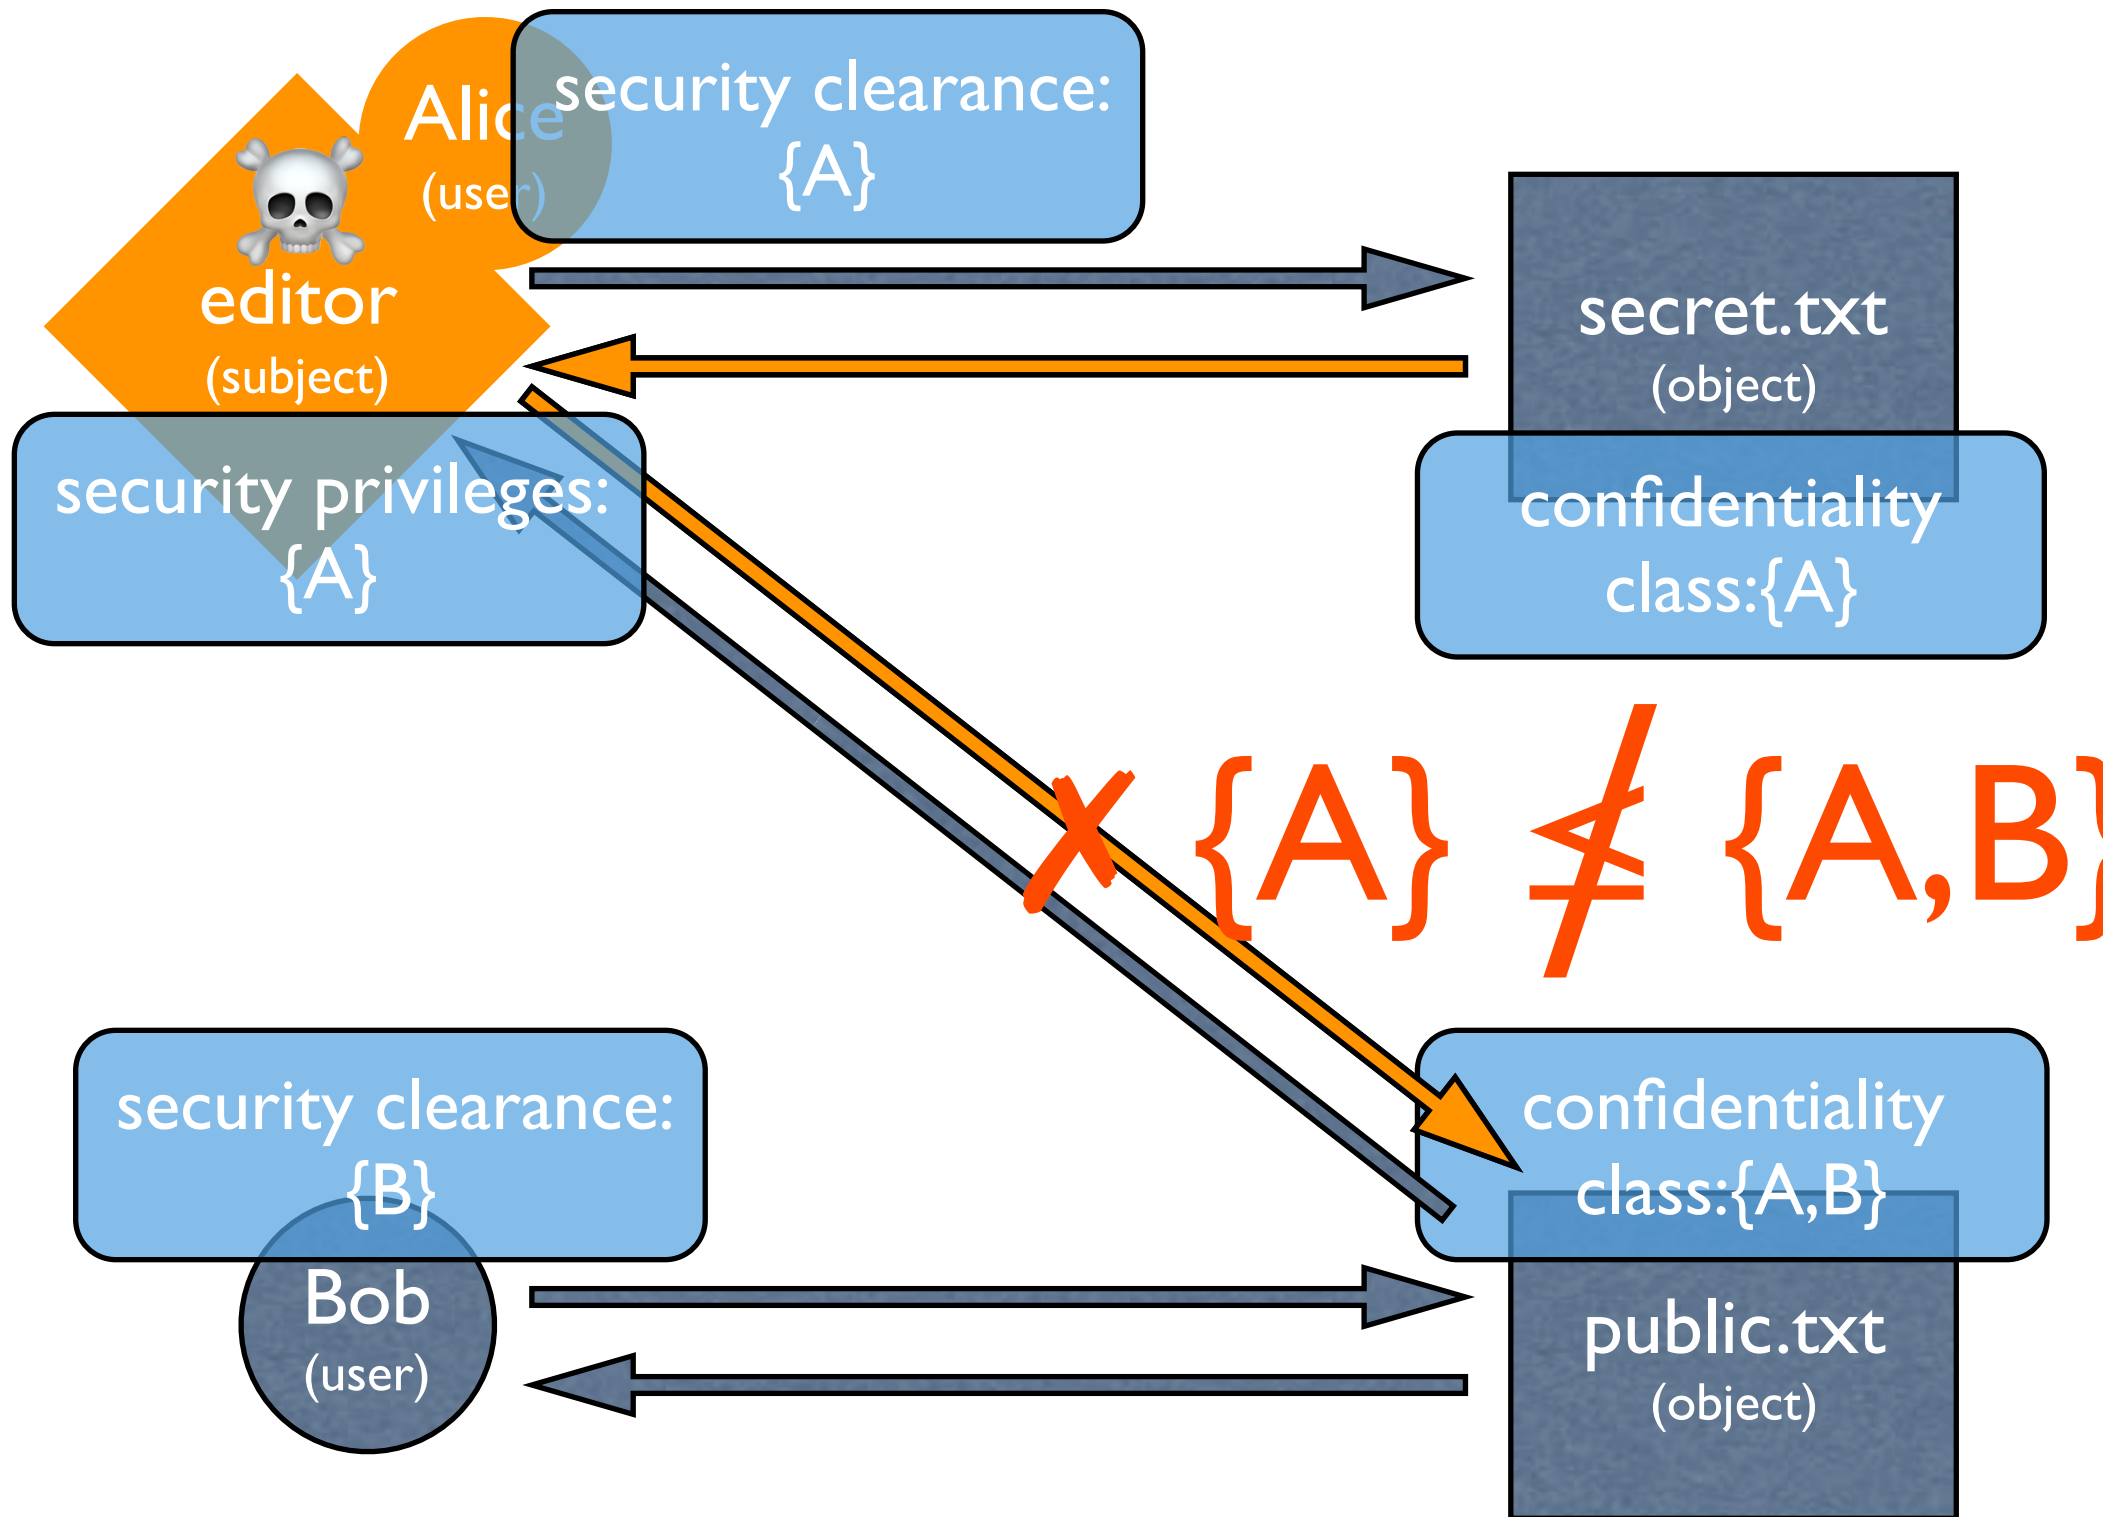
# Example: BLP

- Confidentiality labels can be compared and are assigned to subjects (s) and objects (o)
- BLP mandatory access rules
  - objects can only be read by subjects with the same or higher confidentiality clearance
  - subjects can only write to objects with the same or higher confidentiality level








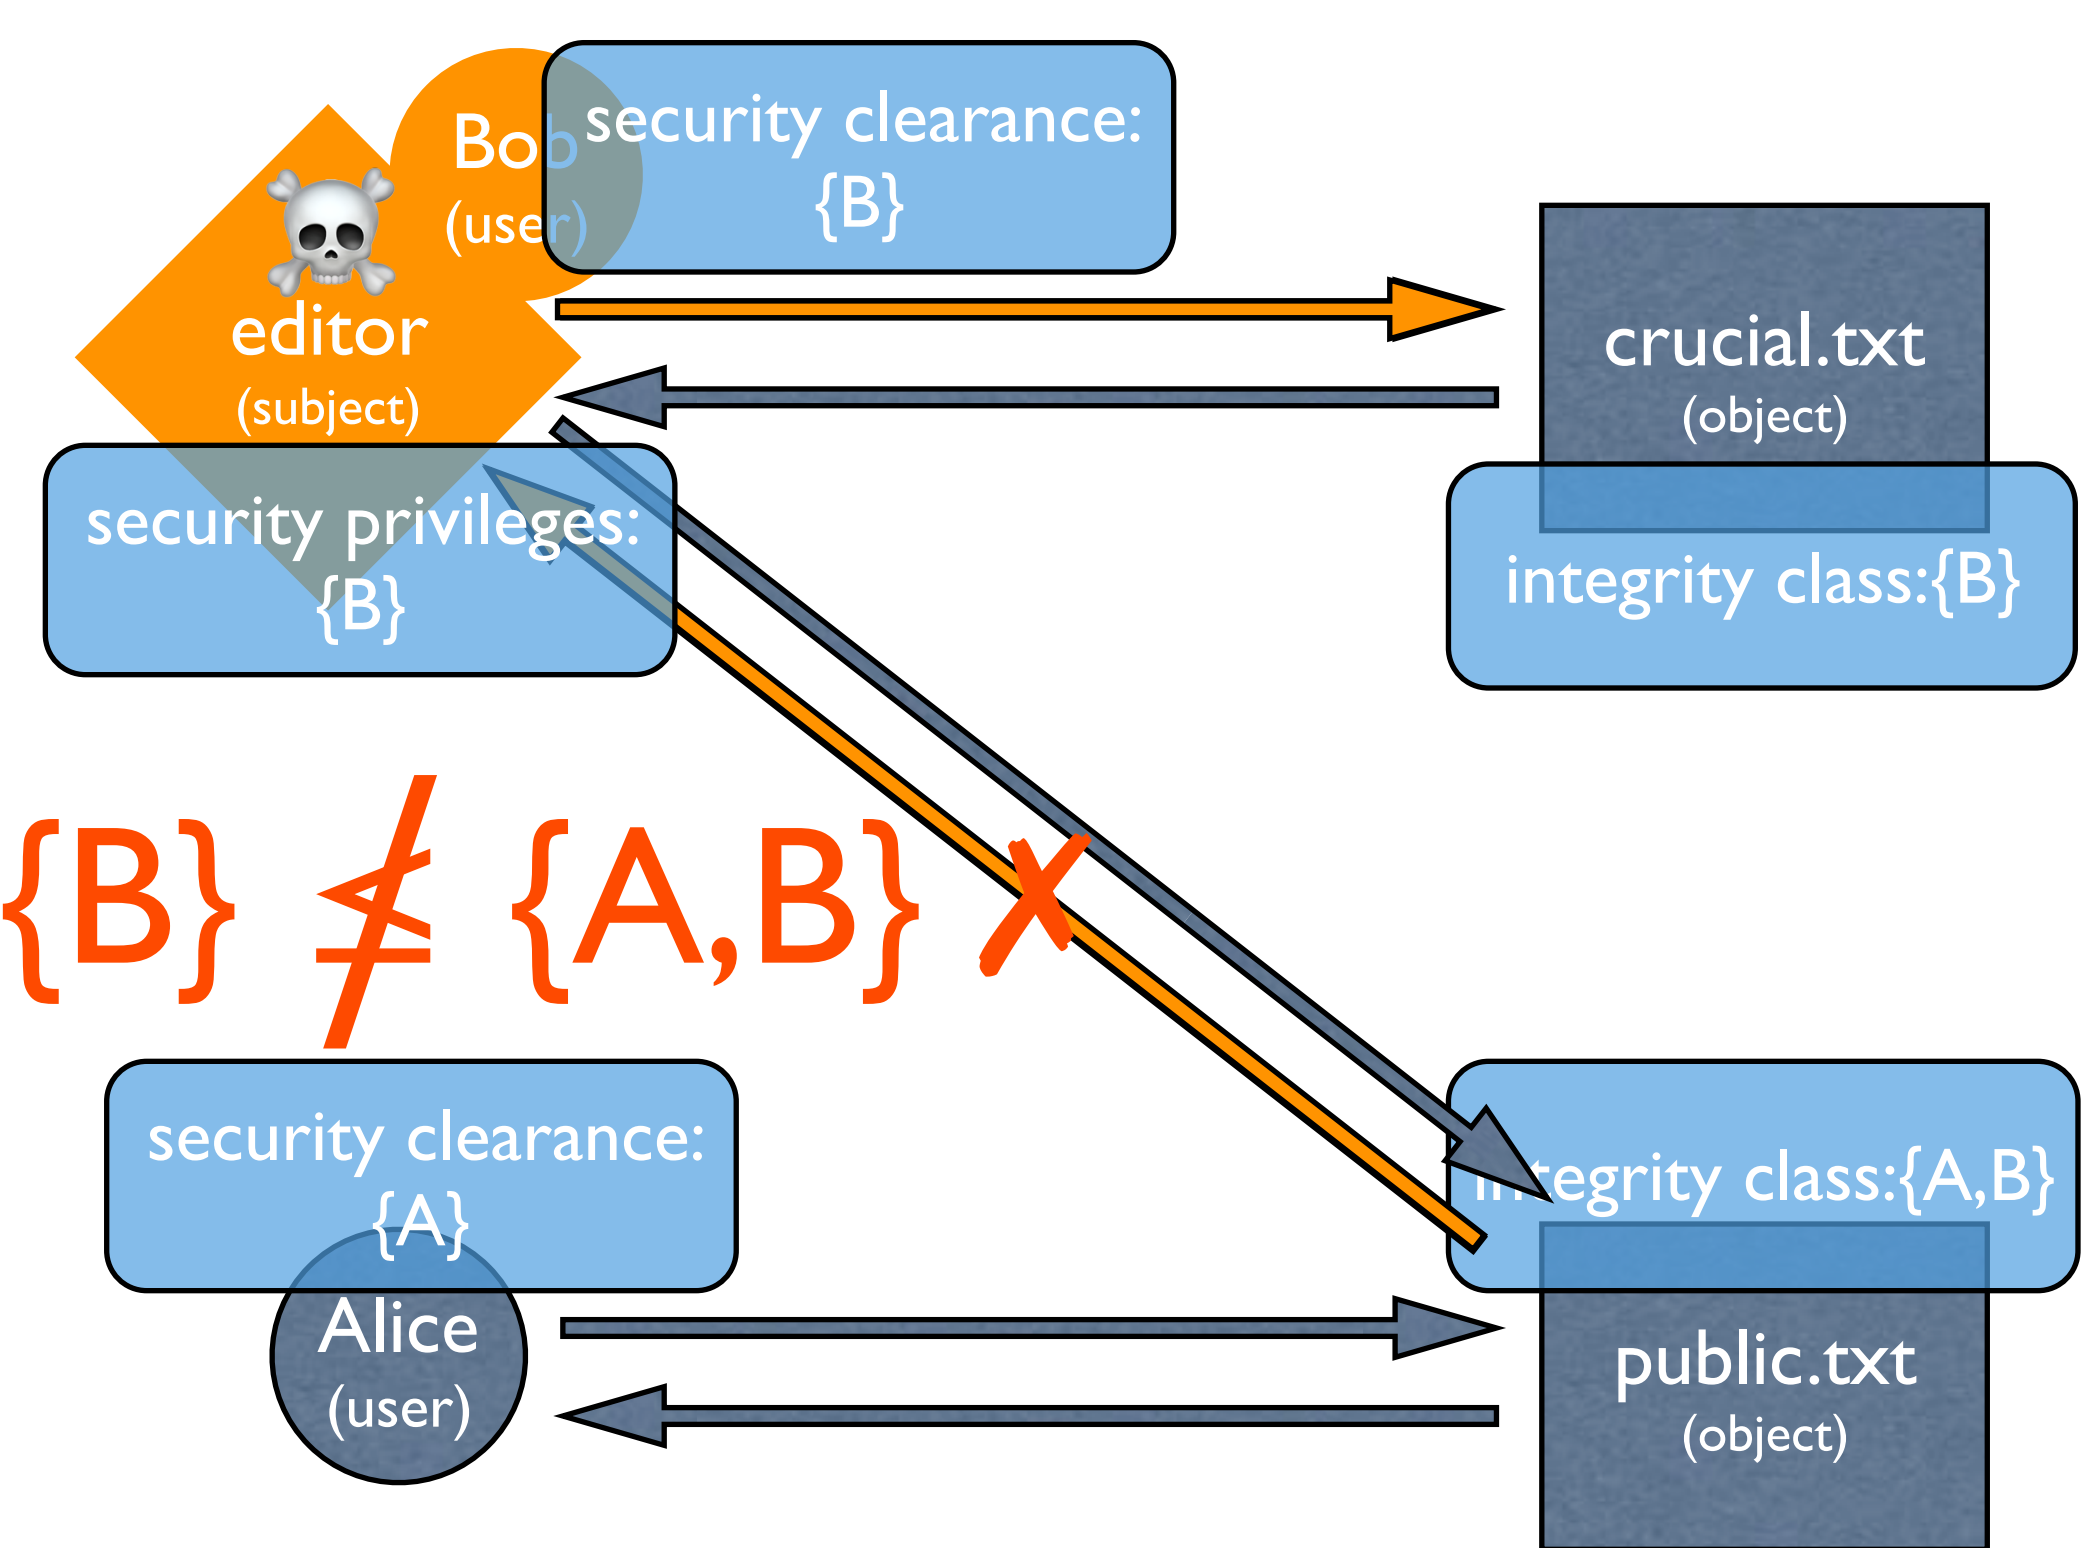


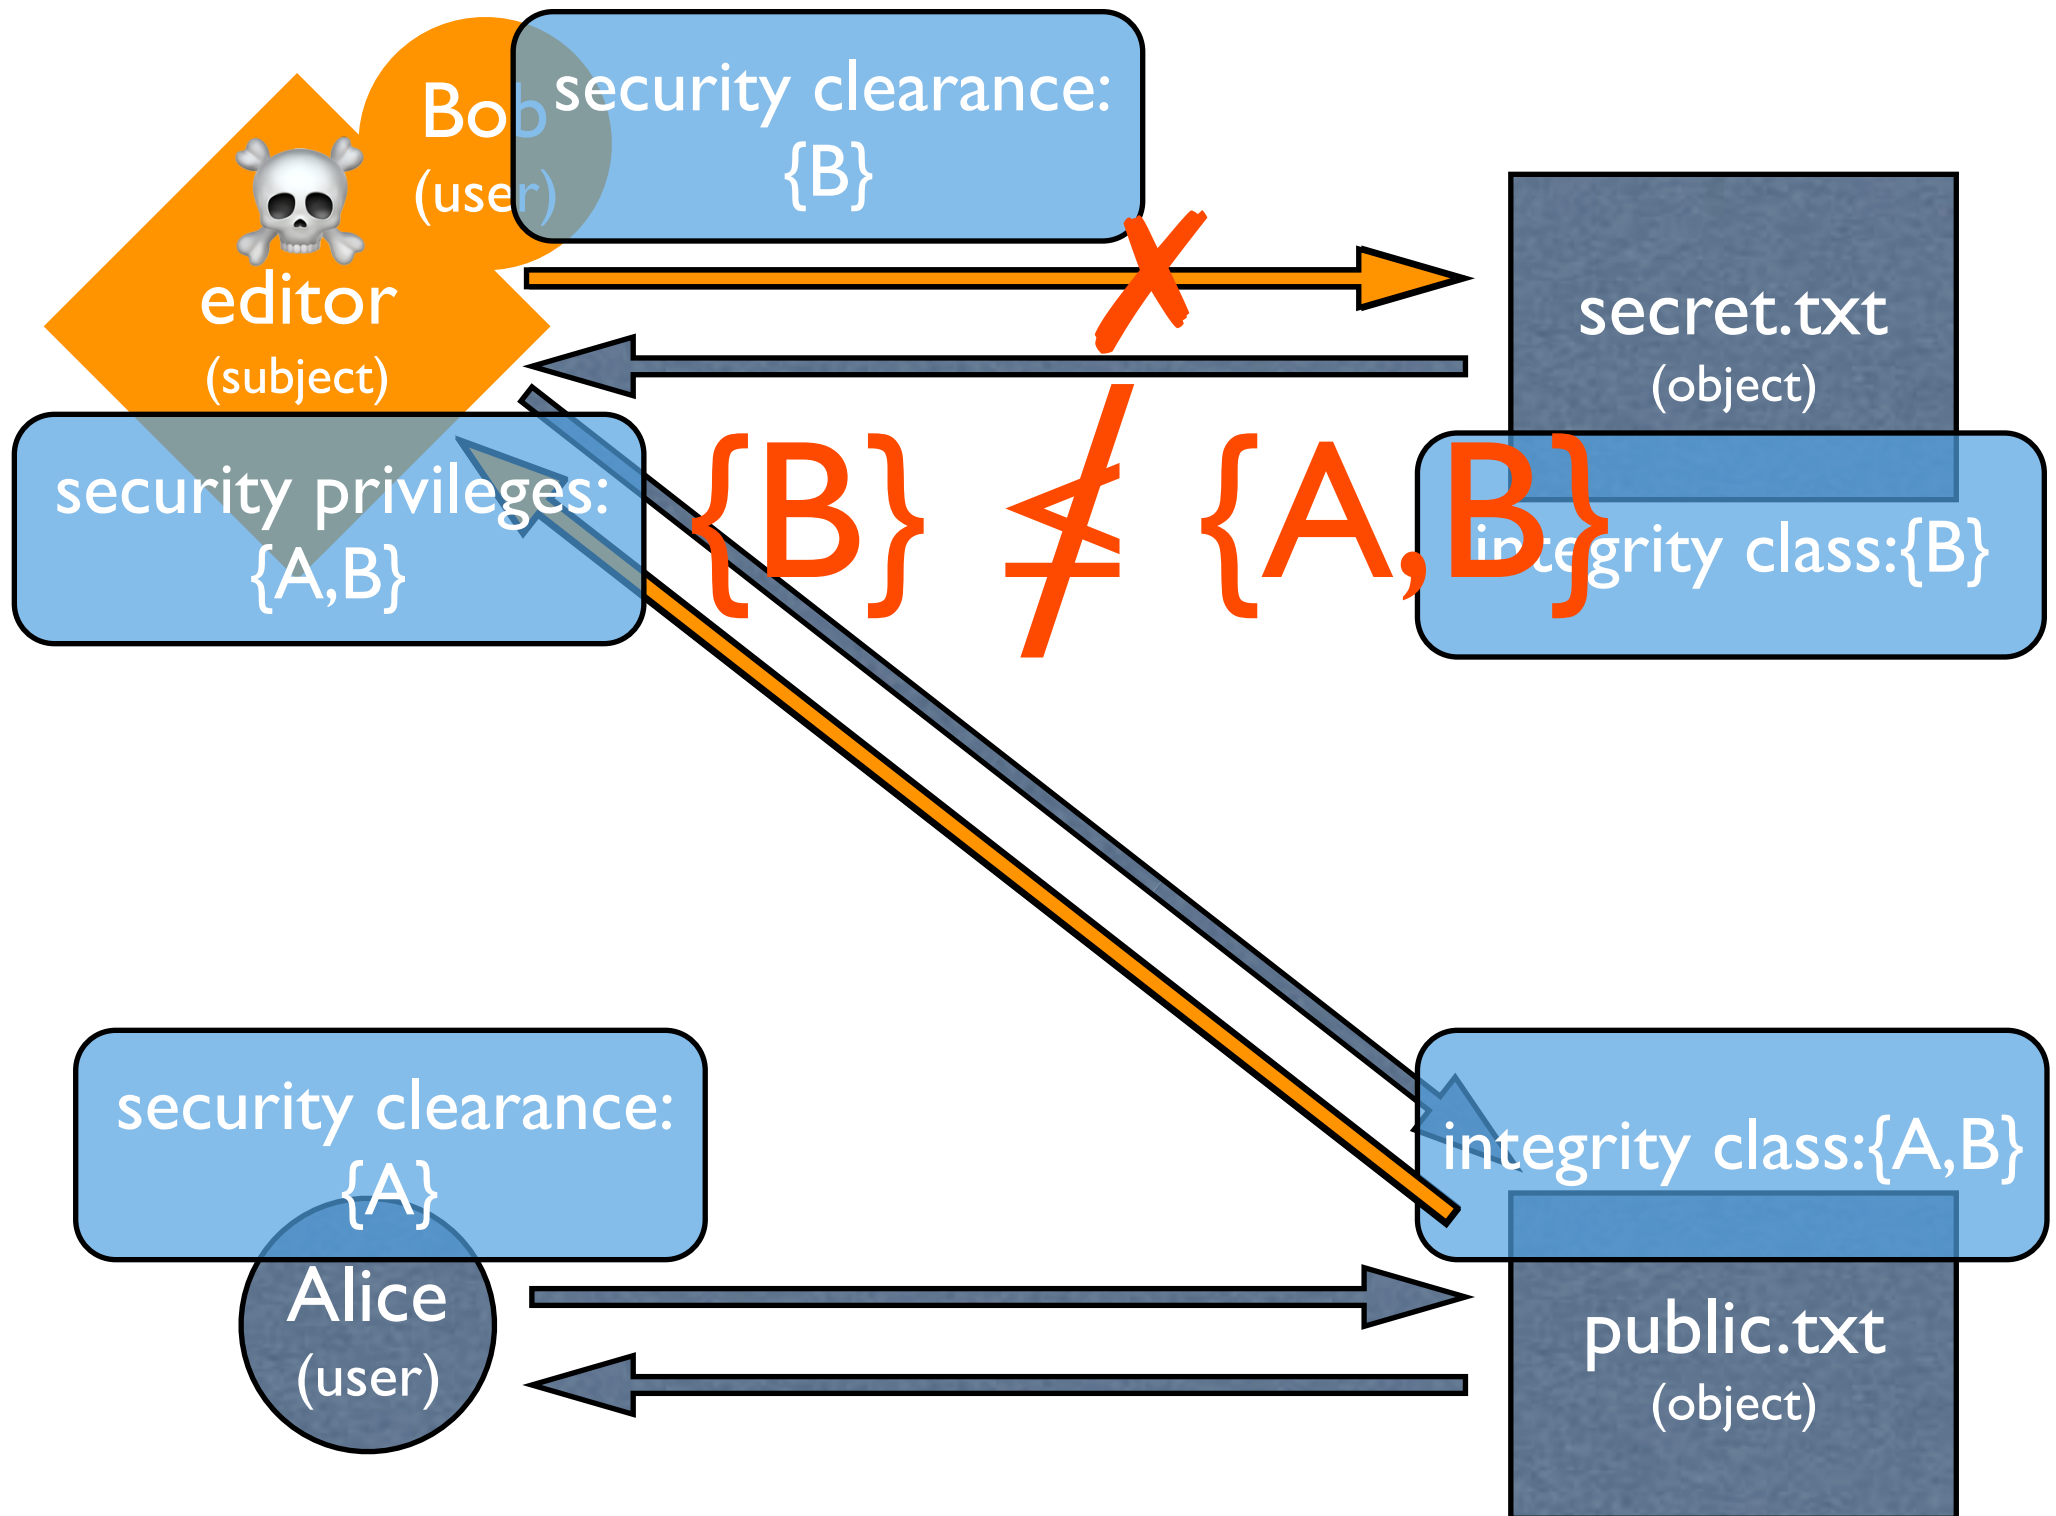
[Biba, 1977]

# Example: Biba

- Integrity labels can be compared and are assigned to subjects (s) and objects (o)
- Biba mandatory access rules
  - subjects can only read objects with the same or higher integrity level
  - objects can only be written by subjects with the same or higher integrity clearance







# Confidentiality + Integrity: a composite model

- Confidentiality and integrity labels form a lattice (product of a confidentiality and an integrity lattice)
- $\lambda$  and  $\omega$  assign confidentiality and integrity labels to subjects (s) and objects (o)
- Mandatory access rules:
  - simple security: s can read o only if  $\lambda(o) \leq \lambda(s)$  and  $\omega(s) \leq \omega(o)$
  - ★-security: s can write to o only if  $\lambda(s) \leq \lambda(o)$  and  $\omega(o) \leq \omega(s)$



editor  
(subject)

Eva  
(user)

security clearance:  
 $\{E\}$

security privileges:  
 $\{E\}$

private.txt  
(object)

confidentiality &  
integrity class:  $\{E\}$

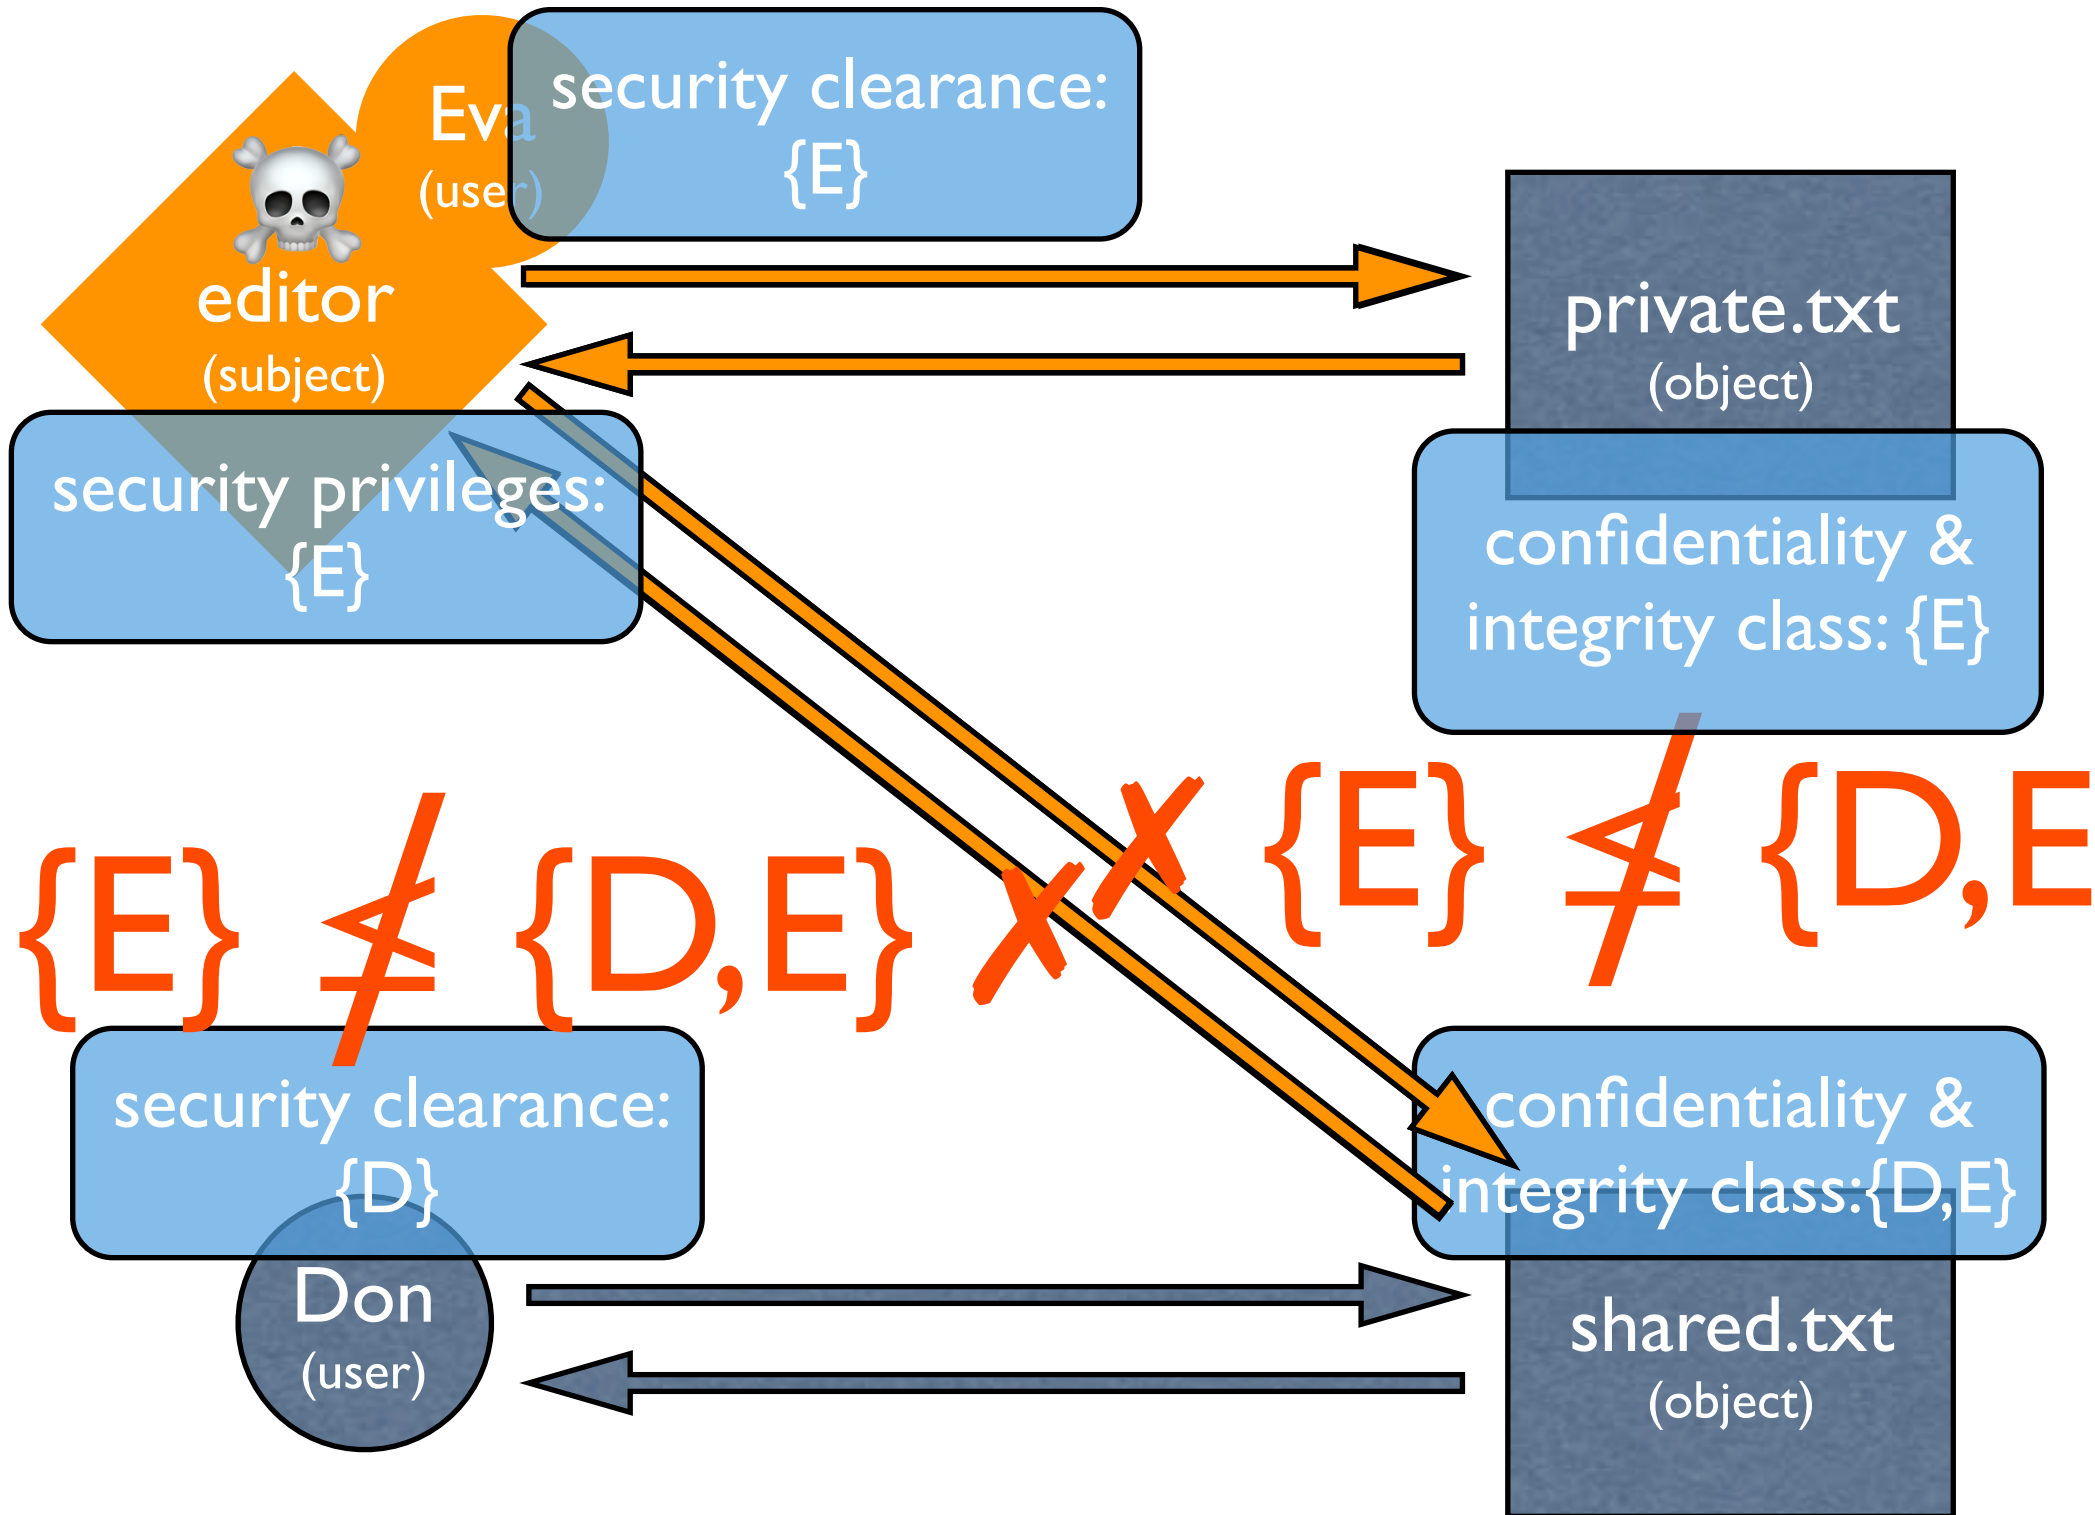
$\{E\} \not\leq \{D, E\}$   $\times \times \{E\} \not\leq \{D, E\}$

security clearance:  
 $\{D\}$

Don  
(user)

confidentiality &  
integrity class:  $\{D, E\}$

shared.txt  
(object)





editor  
(subject)

Eva  
(user)

security clearance:  
 $\{E\}$

security privileges:  
 $\{D, E\}$

$\{E\}$

$\neq$

$\{D, E\}$

$\{E\}$

$\neq$

$\{D, E\}$

security clearance:  
 $\{D\}$

Don  
(user)

confidentiality &  
integrity class:  $\{D, E\}$

shared.txt  
(object)

private.txt  
(object)

confidentiality &  
integrity class:  $\{E\}$



# Example

Authority of the process	secret.txt	crucial.txt	public.txt
Alice	r	-	r
Bob	-	w	w

- Now can you conceive a program that discloses Alice's secret information?  
Or ruins Bob's crucial information?
- What are the disadvantages of these restrictions?



# Covert information flow

- Consider two files:
  - “secret.txt” - owned by Alice, can only be read and written by Alice
  - “public.txt” - can be read and written by everyone
- Can you conceive a program that discloses the secret information to an attacker that can see non-termination?
  - Can Discretionary Access Control prevent it?
  - Can Mandatory Access Control prevent it?

# From Access Control...

- In the absence of knowledge about what a given program does, decisions are necessarily coarse to ensure confidentiality and integrity.
- Discretionary access control cannot ensure end-to-end policies.
- Mandatory access control can do so using very restrictive policies.

# ... to Information Flow Control

“By decoupling the **right to access** information from the **right to disseminate it**, the flow model goes beyond the access matrix model in its ability to specify secure information flow. A practical system needs **both** access and flow control to satisfy all security requirements.”


D. Denning, 1976

- Information flow control focuses on how information is allowed to flow once an access control is granted.

# Class Outline

- Information Flow Security
  - introduction
- Access Control to Information Flow Control
- Encoding and exploiting information flows

# Encoding and exploiting information flows

 “Language-Based Information-Flow Security”, A. Sabelfeld and A. Myers , 2002.

# “no illegal flows” property

- We want to ensure that propagation of information by programs respects information flow policies, i.e. there are no illegal flows
- “no illegal flows” = an attacker cannot infer secret input or affect critical output by inserting inputs into the system and observing its outputs.
- How can we express the property of whether a program respects information flow policies?

# Suppose YOU are the attacker

- You can only read and write to variables of “low level” -- i.e., your level or lower
- For each of the following programs:
  - Can you use the program to uncover “high level” information?
  - Can you use the program to affect “high level” information?

# Following examples

- Language: standard imperative language where information containers are variables.
- Notation:  $\mathbf{X}_L$  denotes that a variable  $x$  has security level  $L$ .
- Information flow policy:

Confidentiality      **secret**  
                                |  
                                **you**

Integrity      **you**  
                                |  
                                **untainted**



# Preserves confidentiality?

- $y_{\text{secret}} := x_{\text{you}}$



Explicit  
leak

- $x_{\text{you}} := y_{\text{secret}}$



- if  $y_{\text{secret}}$  then  $x_{\text{you}} := 0$  else  $x_{\text{you}} := 0$



- if  $y_{\text{secret}}$  then  $x_{\text{you}} := 0$  else  $x_{\text{you}} := 1$



Implicit  
leak

- if  $x_{\text{you}}$  then  $y_{\text{secret}} := 0$  else  $y_{\text{secret}} := 1$



- $x_{\text{you}} := 0$  ; while  $x_{\text{you}} < y_{\text{secret}}$  do  $x_{\text{you}} := x_{\text{you}} + 1$



- $x_{\text{you}} := 1$  ; while  $y_{\text{secret}}$  do  $x_{\text{you}} := 0$














Depends

- while  $x_{\text{you}}$  do skip ;  $y_{\text{secret}} := 0$



# Preserves integrity?

- $y_{\text{untainted}} := x_{\text{you}}$   
- $x_{\text{you}} := y_{\text{untainted}}$  
- if  $y_{\text{untainted}}$  then  $x_{\text{you}} := 0$  else  $x_{\text{you}} := 0$  
- if  $y_{\text{untainted}}$  then  $x_{\text{you}} := 0$  else  $x_{\text{you}} := 1$  
- if  $x_{\text{you}}$  then  $y_{\text{untainted}} := 0$  else  $y_{\text{untainted}} := 1$   
- $y_{\text{untainted}} := 0$  ; while  $x_{\text{you}} > y_{\text{untainted}}$  do  $y_{\text{untainted}} := y_{\text{untainted}} + 1$  
- $x_{\text{you}} := 1$  ; while  $y_{\text{secret}}$  do  $x_{\text{you}} := 0$  
- while  $x_{\text{you}}$  do skip ;  $y_{\text{untainted}} := 0$   

# What is the power of the attacker?


- Can the attacker observe:
  - That a program does not terminate?
  - Intermediate steps of the computation?
  - The possibility of producing certain states?
  - The likelihood of producing certain states?

# Challenge for next class

- Secure program = the program does not encode illegal information flows
- Can you formulate a security property that defines when a program is secure?

# Tips

- Suppose you are the attacker. You are given a black box program, and you can control the low inputs of the program, and read its low outputs.
- How many executions would you need to find out if the program leaks information?
- What inputs would you give it in order to find out whether the program leaks information?



- Which of the following experiments allows you to conclude whether the program encodes leaks?
  - Input:  $y_L=0$ , Output:  $y_L=0$
  - Input:  $y_L=0$ , Output:  $y_L=1$
  - Input:  $y_L=0$ , Output:  $y_L=0$  and Input:  $y_L=0$ , Output:  $y_L=0$
  - Input:  $y_L=0$ , Output:  $y_L=0$  and Input:  $y_L=1$ , Output:  $y_L=1$
  - Input:  $y_L=0$ , Output:  $y_L=0$  and Input:  $y_L=0$ , Output:  $y_L=1$

# Conclusion

- Access Control mechanisms do not have the appropriate refinement for enforcing Information Flow policies, as they do not take into account program behaviors.
- In order to analyze information flows in a given program, it is necessary to consider its possible behaviors.
- Next class: Definition of Security Properties

